

ORACLE



# ZGC

## The Next Generation Low-Latency Garbage Collector

**Per Liden (@perliden)**

Consulting Member of Technical Staff  
Java Platform Group, Oracle



# Agenda

- 1 What is ZGC?
- 2 The Design of ZGC
- 3 Performance
- 4 Using ZGC
- 5 Roadmap

# Safe harbor statement

The following is intended to outline our general product direction. It is intended for information purposes only, and may not be incorporated into any contract. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions. The development, release, timing, and pricing of any features or functionality described for Oracle's products may change and remains at the sole discretion of Oracle Corporation.





# A Scalable Low-Latency Garbage Collector

# Goals

Max GC pause time

**10<sub>ms</sub>**

Multi-terabyte heaps

**TB**

Max throughput reduction

**15%**

# Goals



Easy to tune!

# GC Landscape

Oracle supported garbage collectors



GC	Optimized For
Serial	Memory Footprint
Parallel	Throughput
G1	Throughput/Latency Balance
ZGC	Low Latency

# ZGC at a Glance



Concurrent  
Tracing  
Compacting  
Single generation

Region-based  
NUMA-aware  
Load barriers  
Colored pointers





ZGC pause times do not increase  
with the heap or live-set size



ZGC pause times do increase  
with the root-set size

(Number of Java Threads)

# What's Concurrent?



	Serial	Parallel	G1	CMS	ZGC
Marking					
Compaction					
Reference Processing					
Relocation Set Selection					
StringTable Cleaning					
JNI WeakRef Cleaning					
JNI GlobalRefs Scanning					
Class Unloading					
Thread Stack Scanning					

# What's Concurrent?



	Serial	Parallel	G1	CMS	ZGC
Marking	-	-			
Compaction	-	-			
Reference Processing	-	-			
Relocation Set Selection	-	-			
StringTable Cleaning	-	-			
JNI WeakRef Cleaning	-	-			
JNI GlobalRefs Scanning	-	-			
Class Unloading	-	-			
Thread Stack Scanning	-	-			

# What's Concurrent?



	Serial	Parallel	G1	CMS	ZGC
Marking	-	-	✓	✓	
Compaction	-	-	-	-	
Reference Processing	-	-	-	-	
Relocation Set Selection	-	-	-	-	
StringTable Cleaning	-	-	-	-	
JNI WeakRef Cleaning	-	-	-	-	
JNI GlobalRefs Scanning	-	-	-	-	
Class Unloading	-	-	-	-	
Thread Stack Scanning	-	-	-	-	

\*) Old Gen only

# What's Concurrent?

	Serial	Parallel	G1	CMS	ZGC
Marking	-	-	✓ *	✓ *	
Compaction	-	-	-	-	
Reference Processing	-	-	-	-	
Relocation Set Selection	-	-	-	-	
StringTable Cleaning	-	-	-	-	
JNI WeakRef Cleaning	-	-	-	-	
JNI GlobalRefs Scanning	-	-	-	-	
Class Unloading	-	-	-	-	
Thread Stack Scanning	-	-	-	-	

# What's Concurrent?

\*) Old Gen only  
 \*\*) Not yet

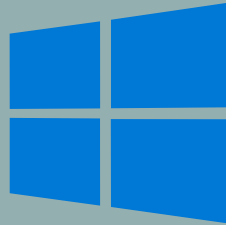
	Serial	Parallel	G1	CMS	ZGC
Marking	-	-	✓ *	✓ *	✓
Compaction	-	-	-	-	✓
Reference Processing	-	-	-	-	✓
Relocation Set Selection	-	-	-	-	✓
StringTable Cleaning	-	-	-	-	✓
JNI WeakRef Cleaning	-	-	-	-	✓
JNI GlobalRefs Scanning	-	-	-	-	✓
Class Unloading	-	-	-	-	✓
Thread Stack Scanning	-	-	-	-	- **

# Supported on All Commonly Used Platforms



Linux

x86 (64-bit)  
Arm (64-bit)



Windows

x86 (64-bit)

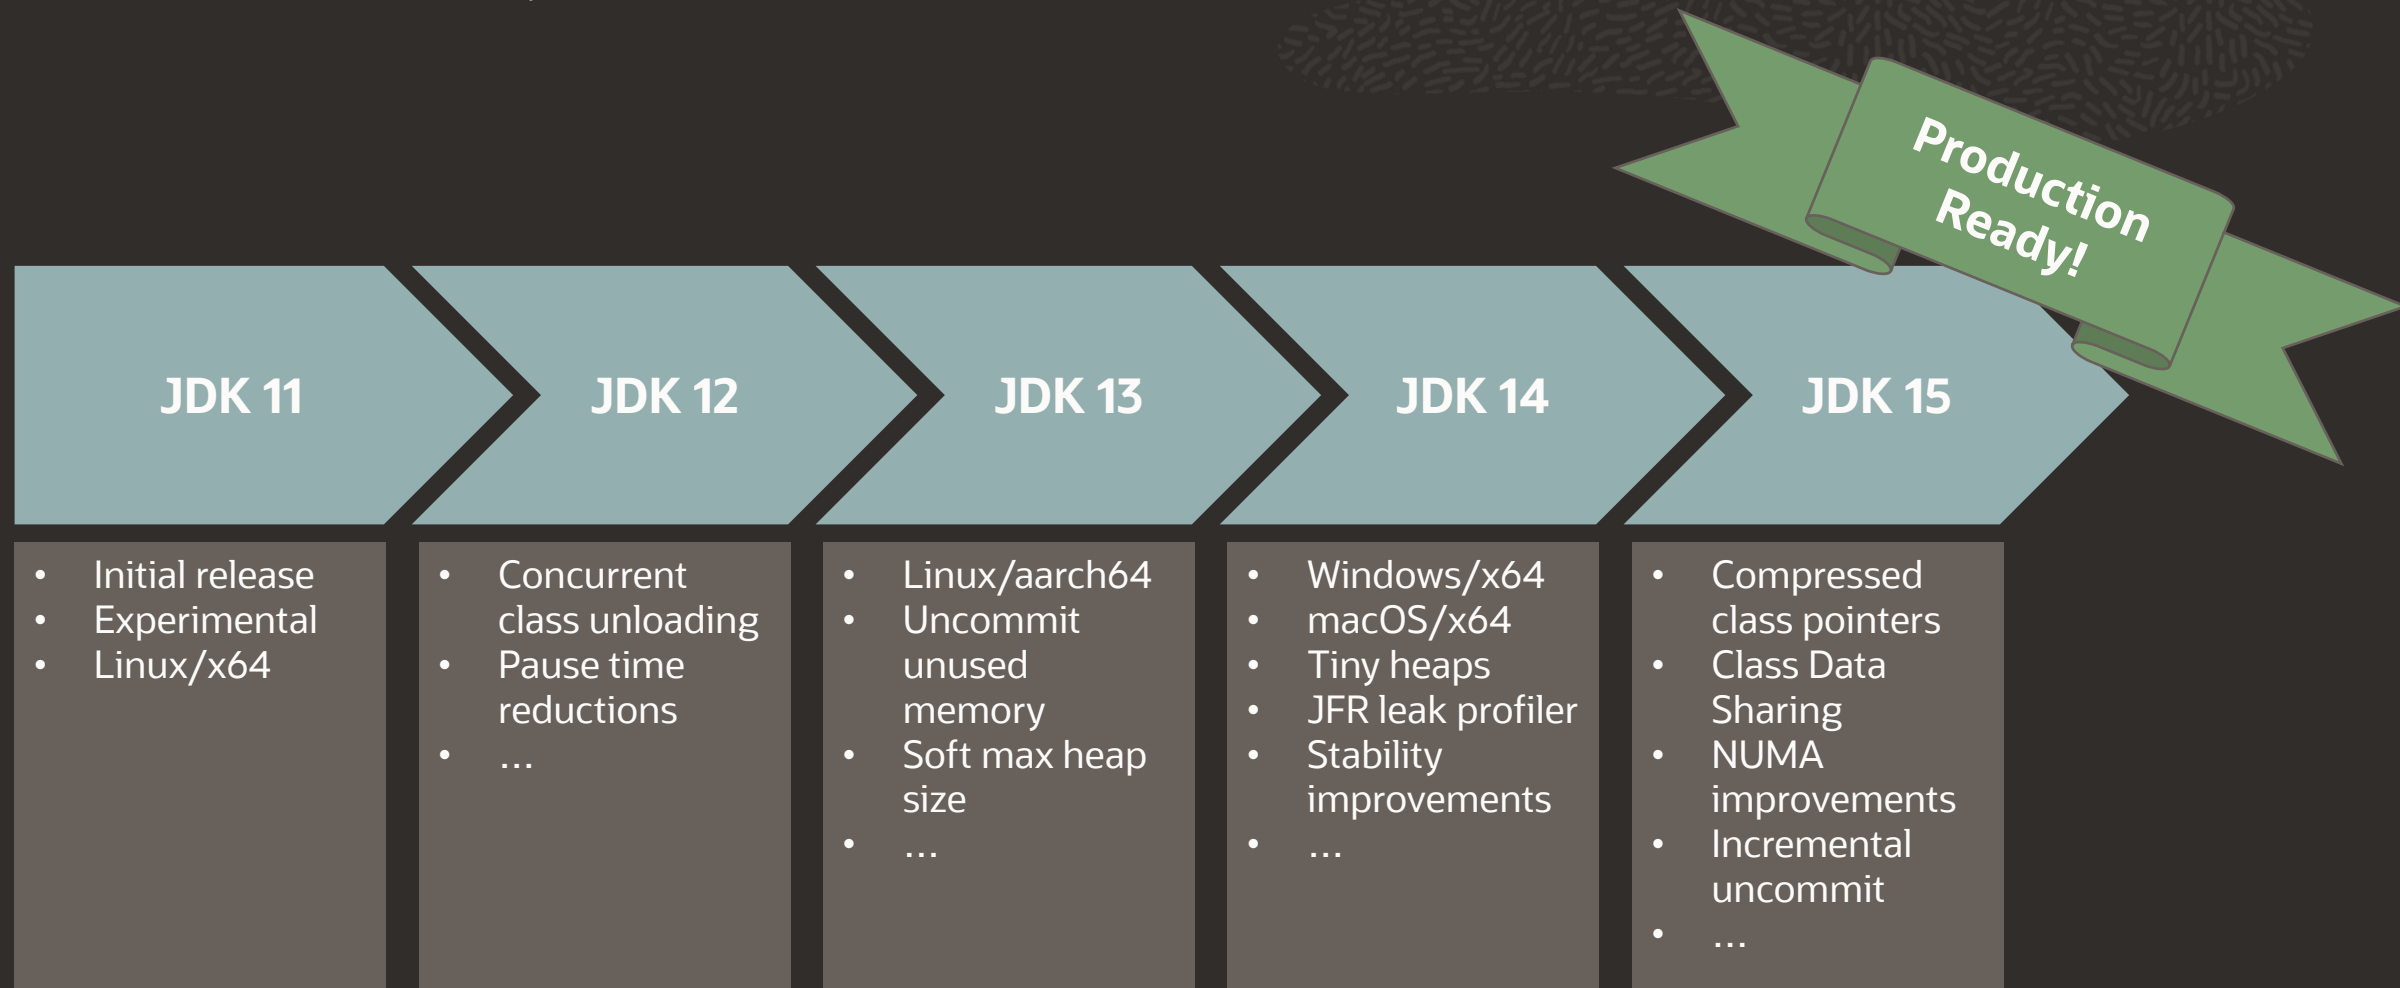


macOS

x86 (64-bit)



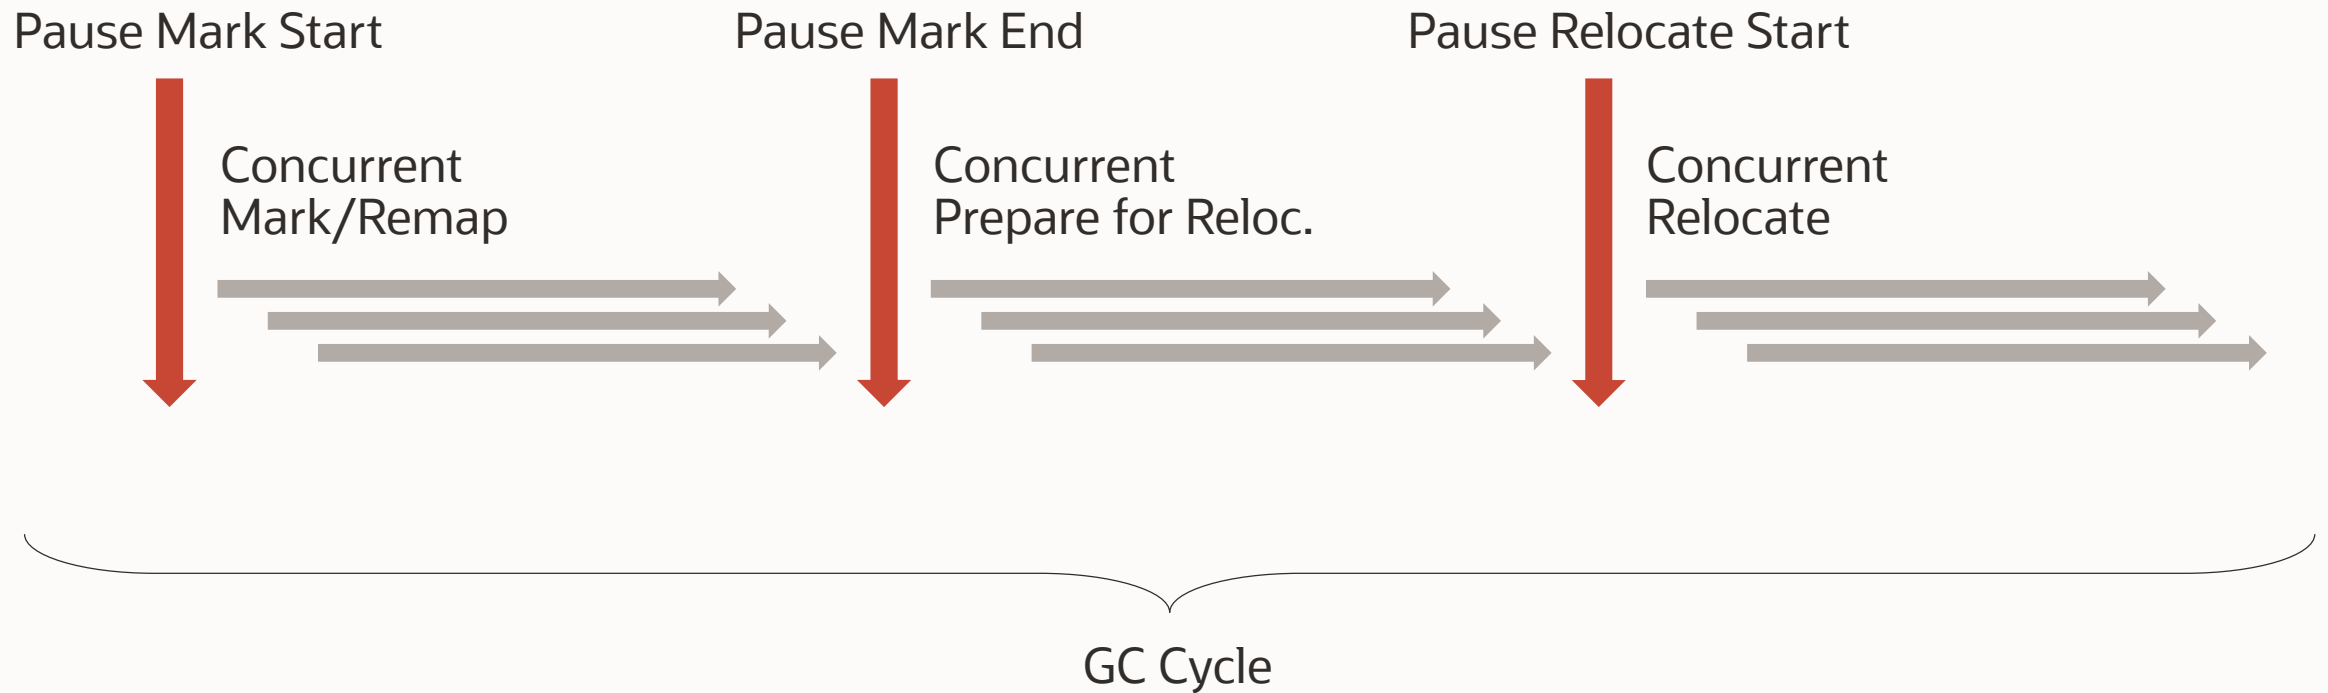
# Production Ready in JDK 15



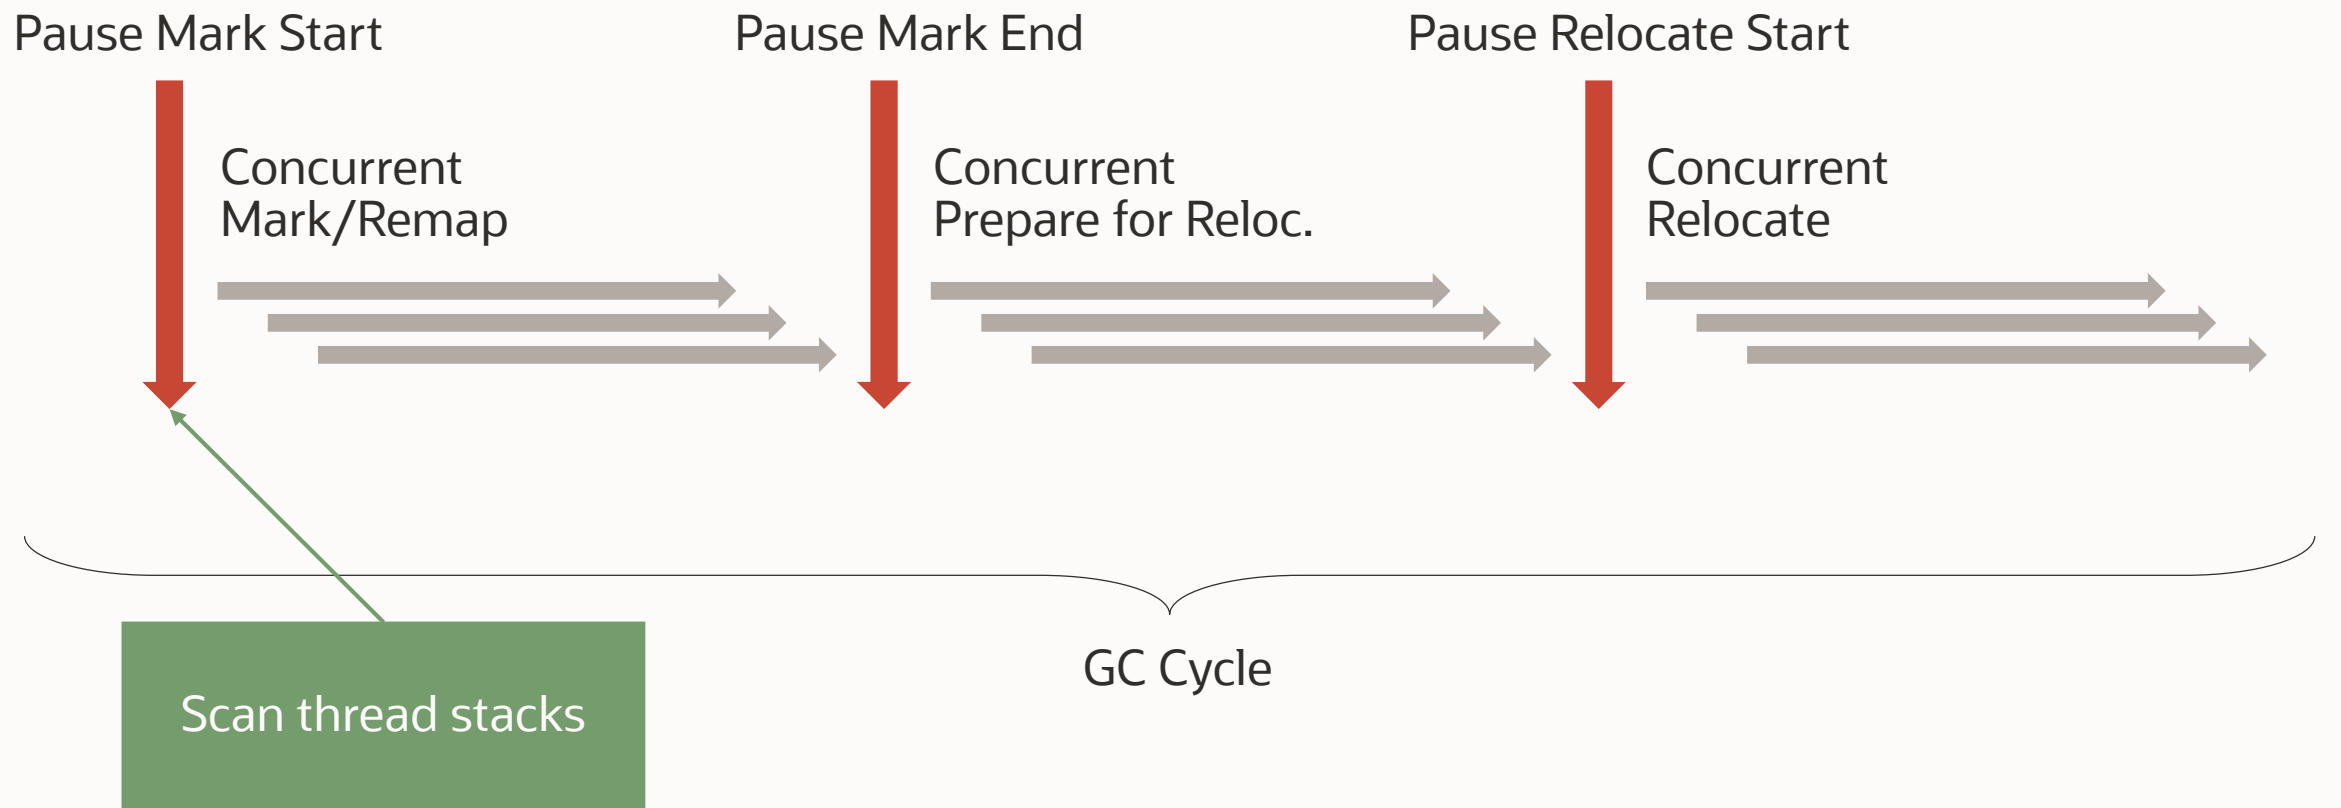


# The Design of ZGC

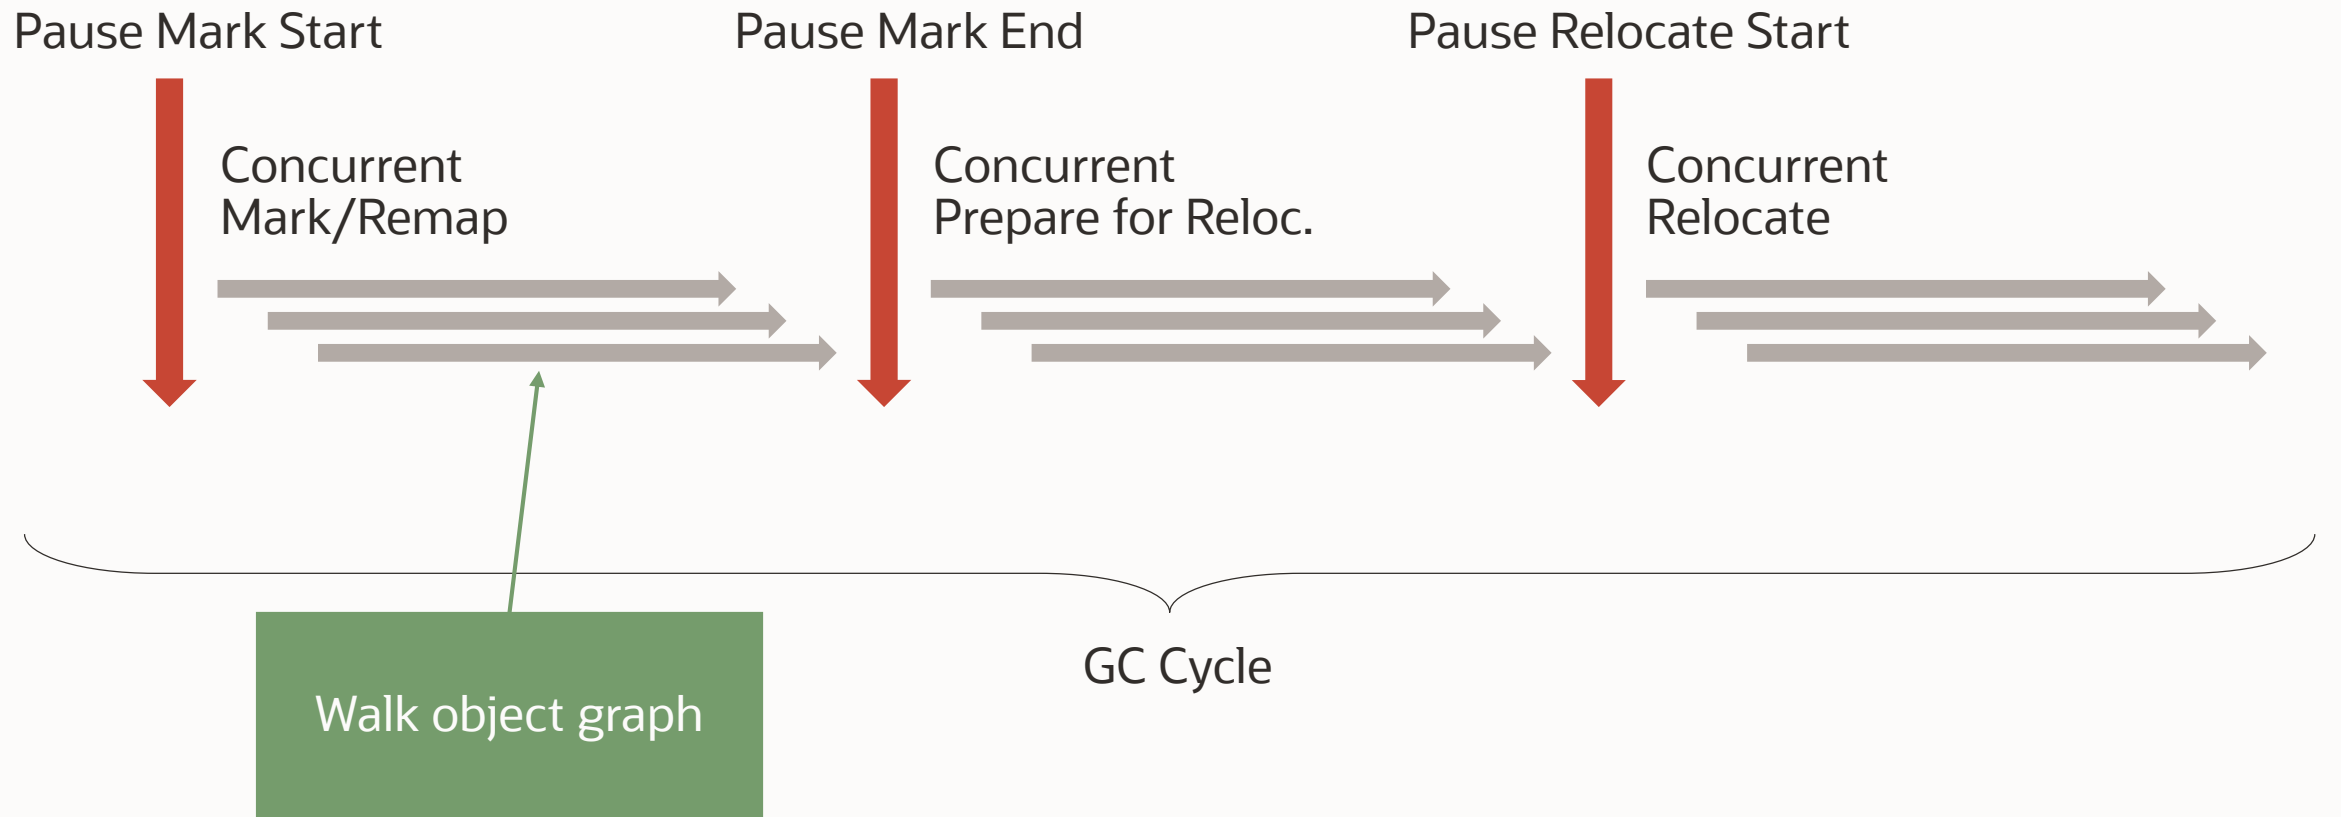
# ZGC Phases



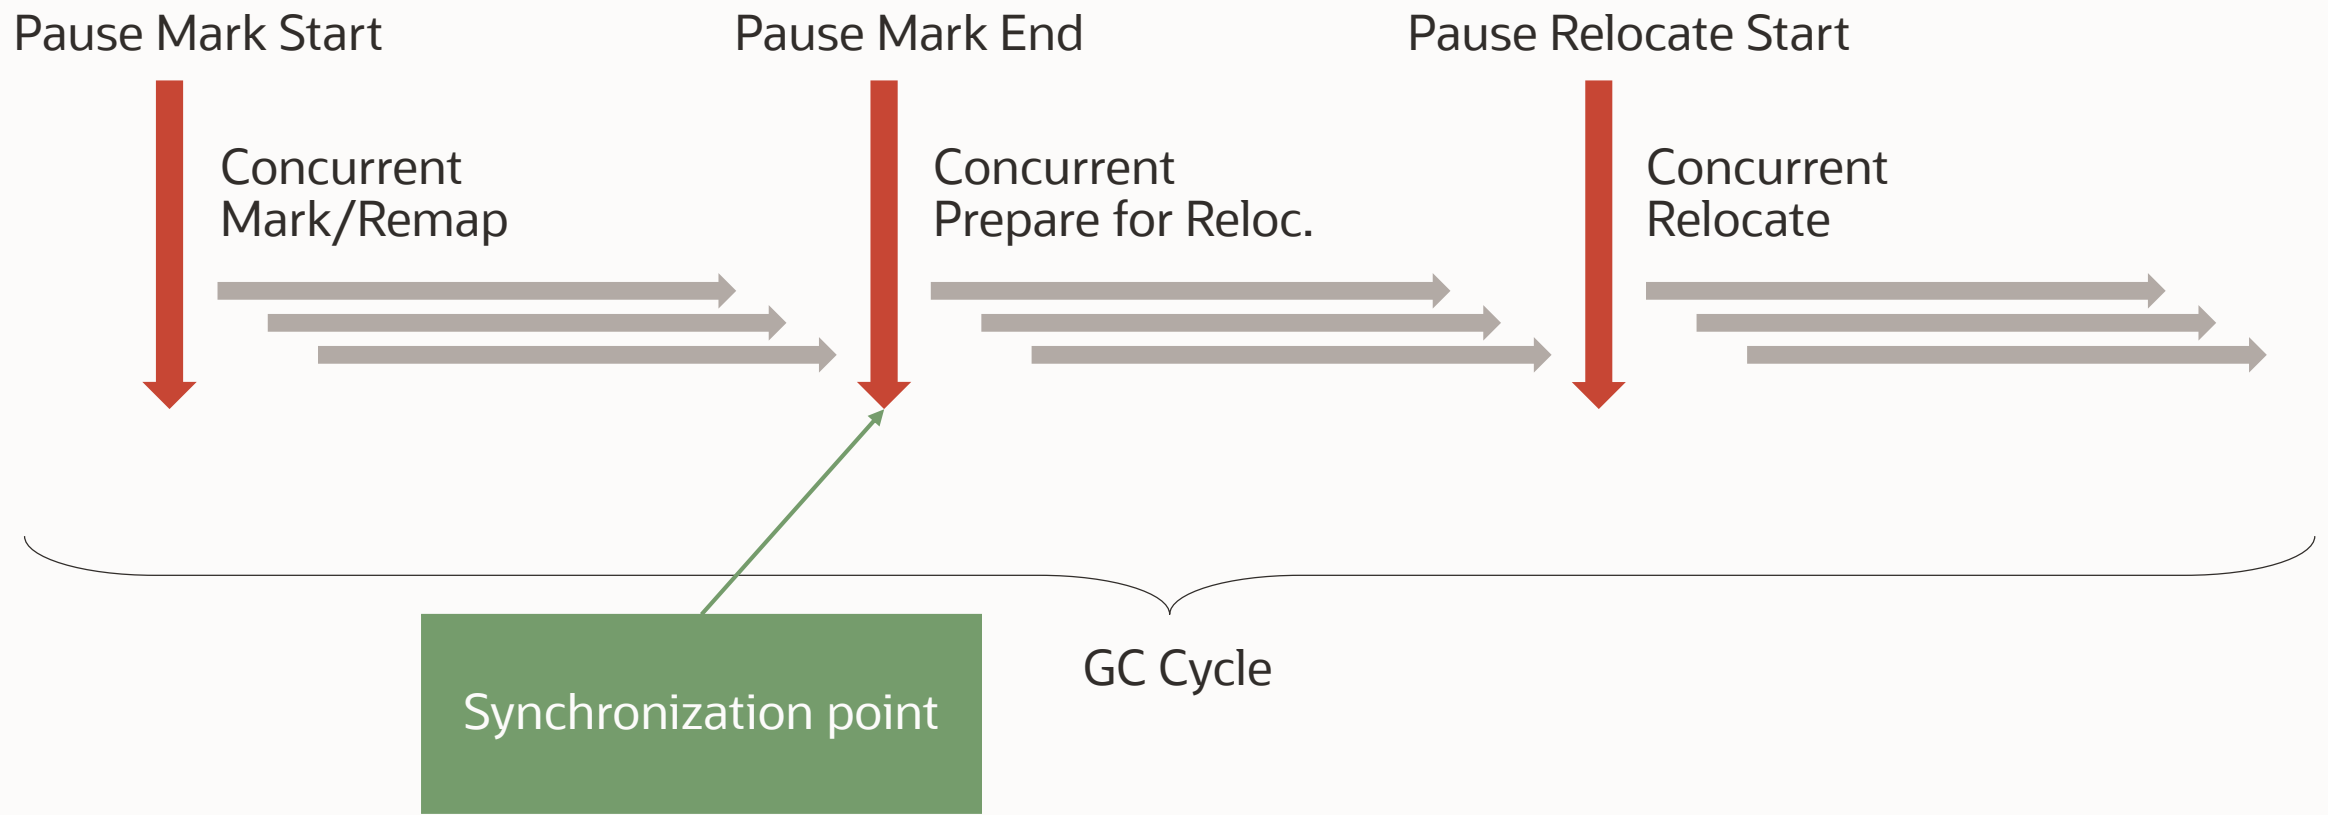
# ZGC Phases



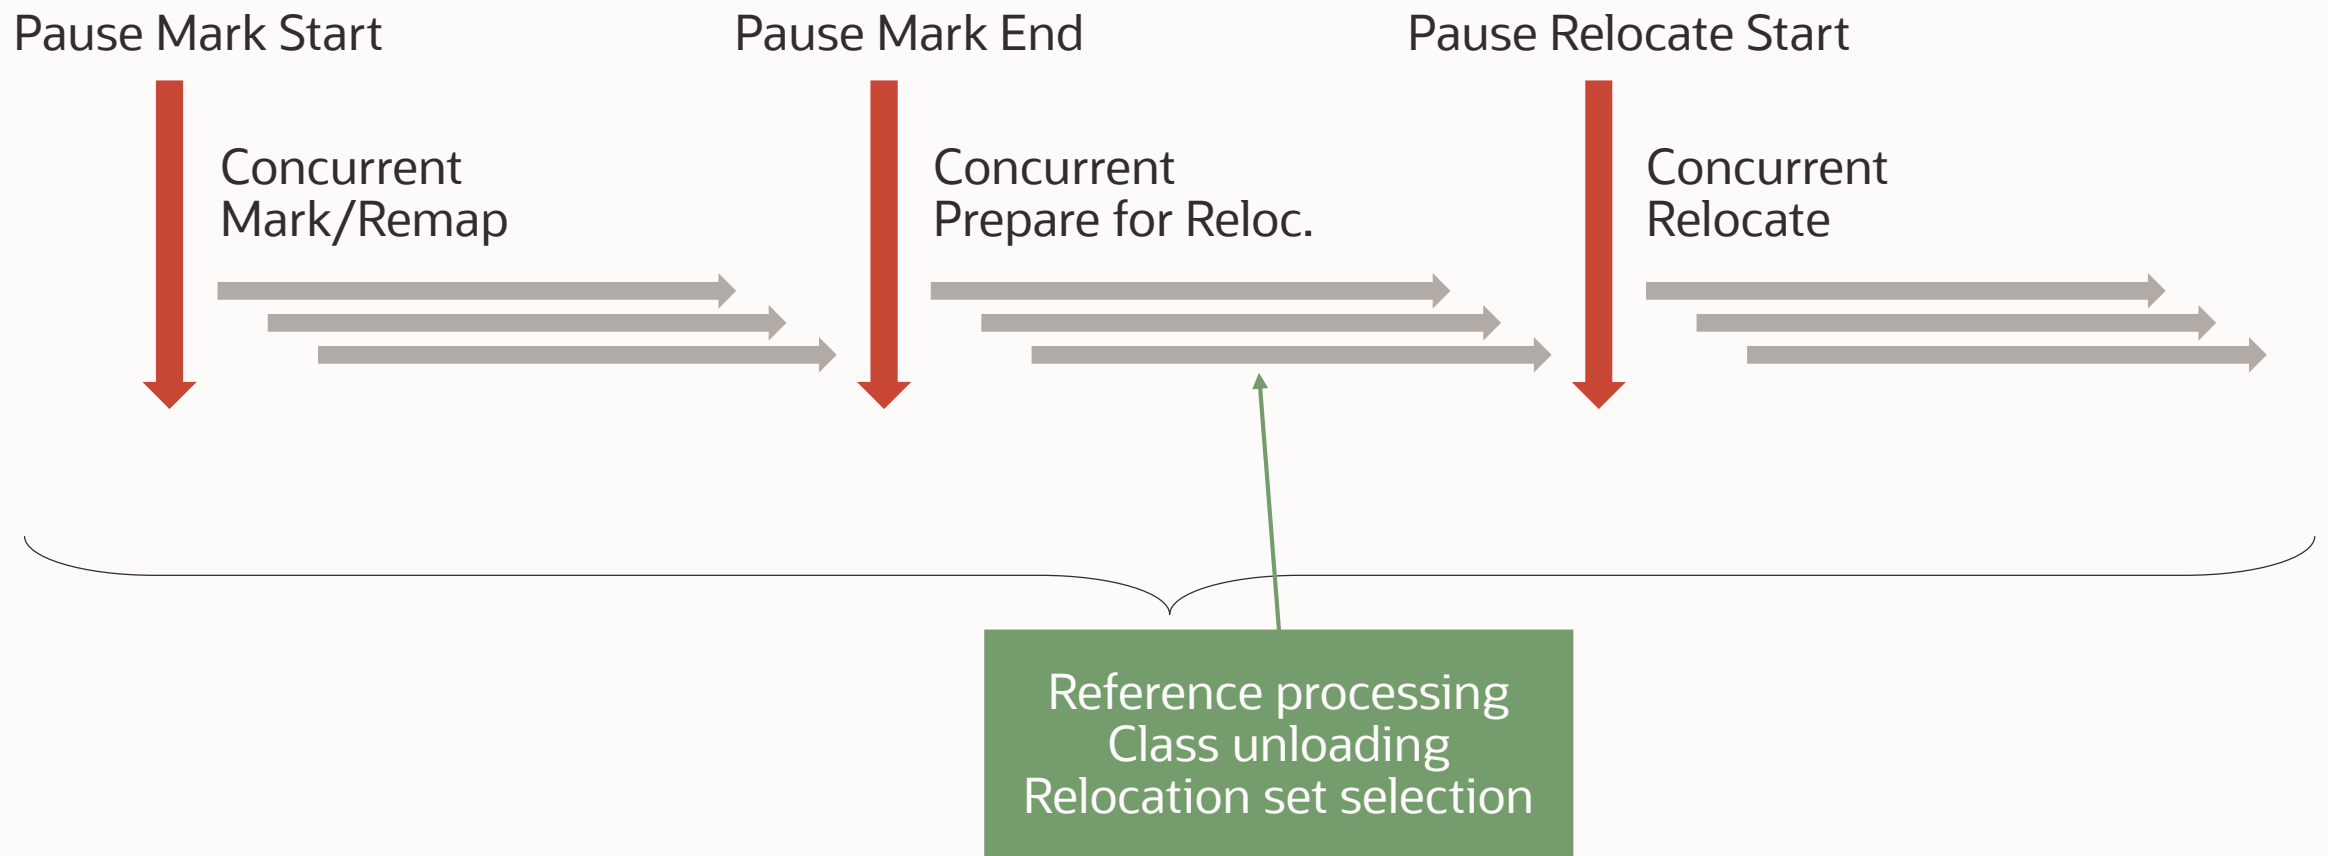
# ZGC Phases



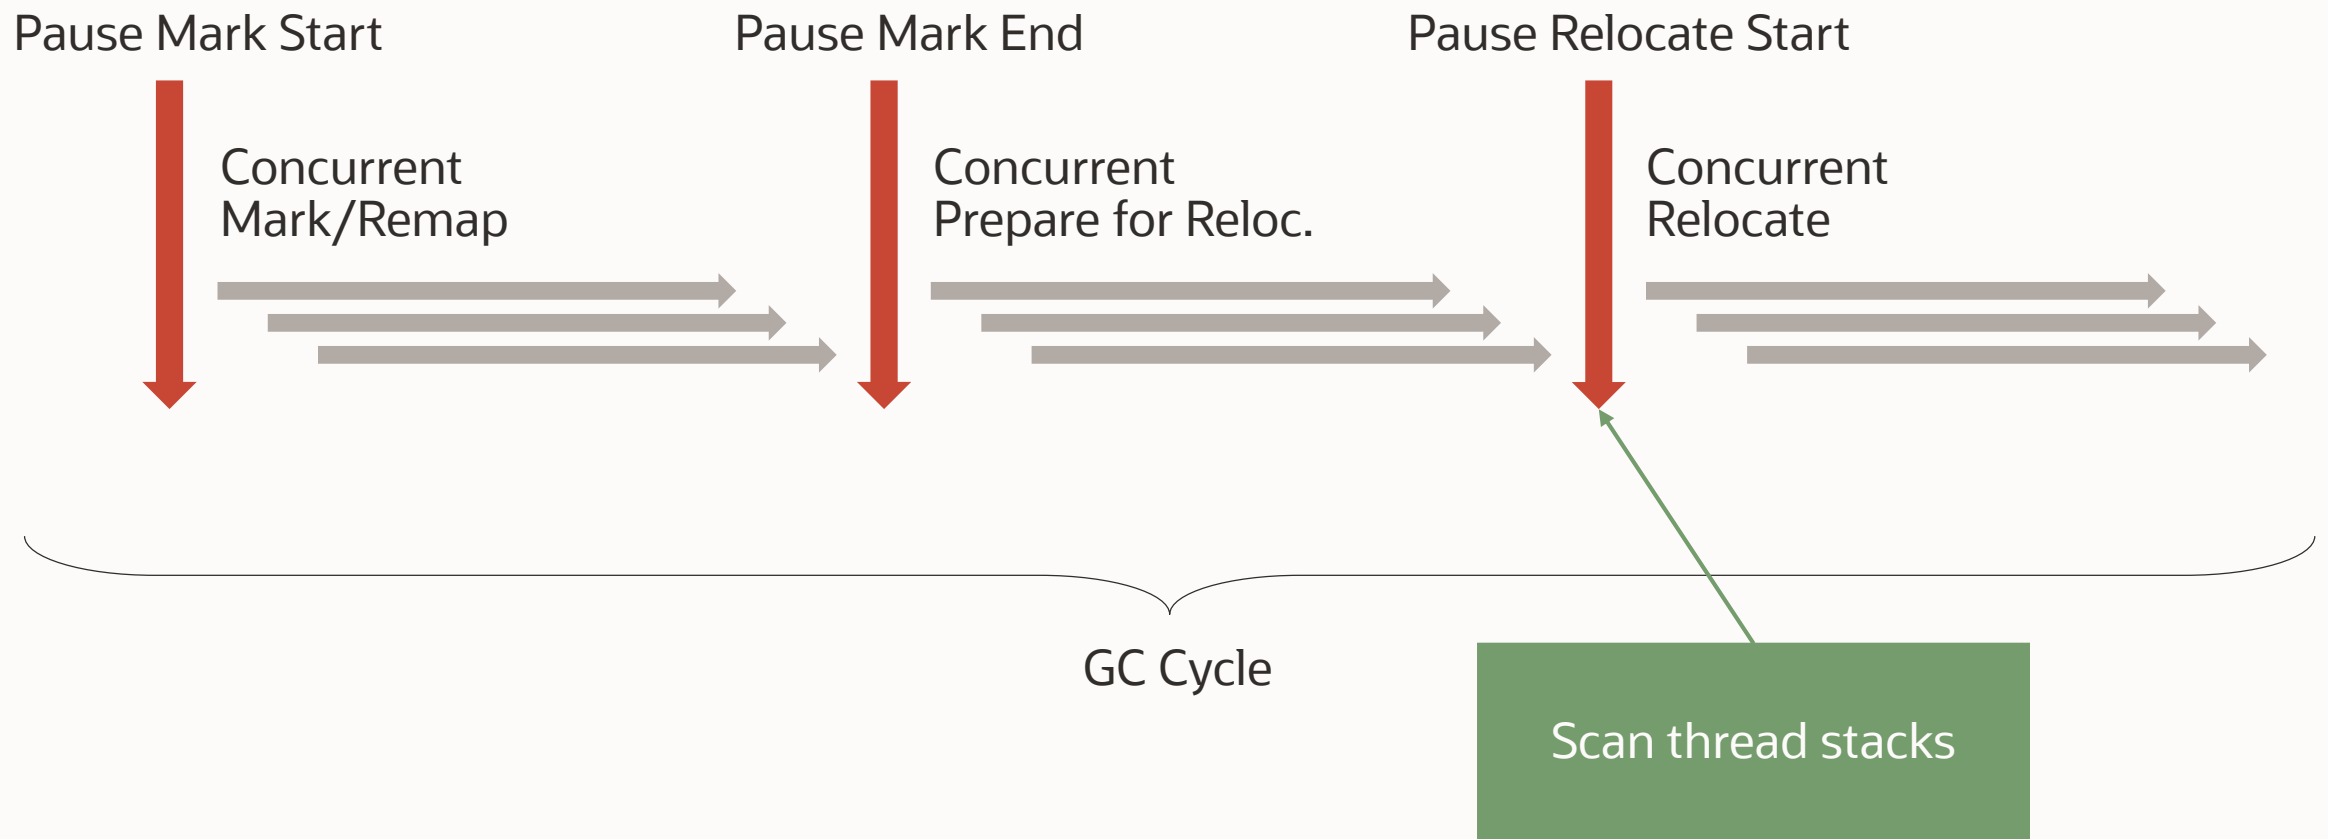
# ZGC Phases



# ZGC Phases

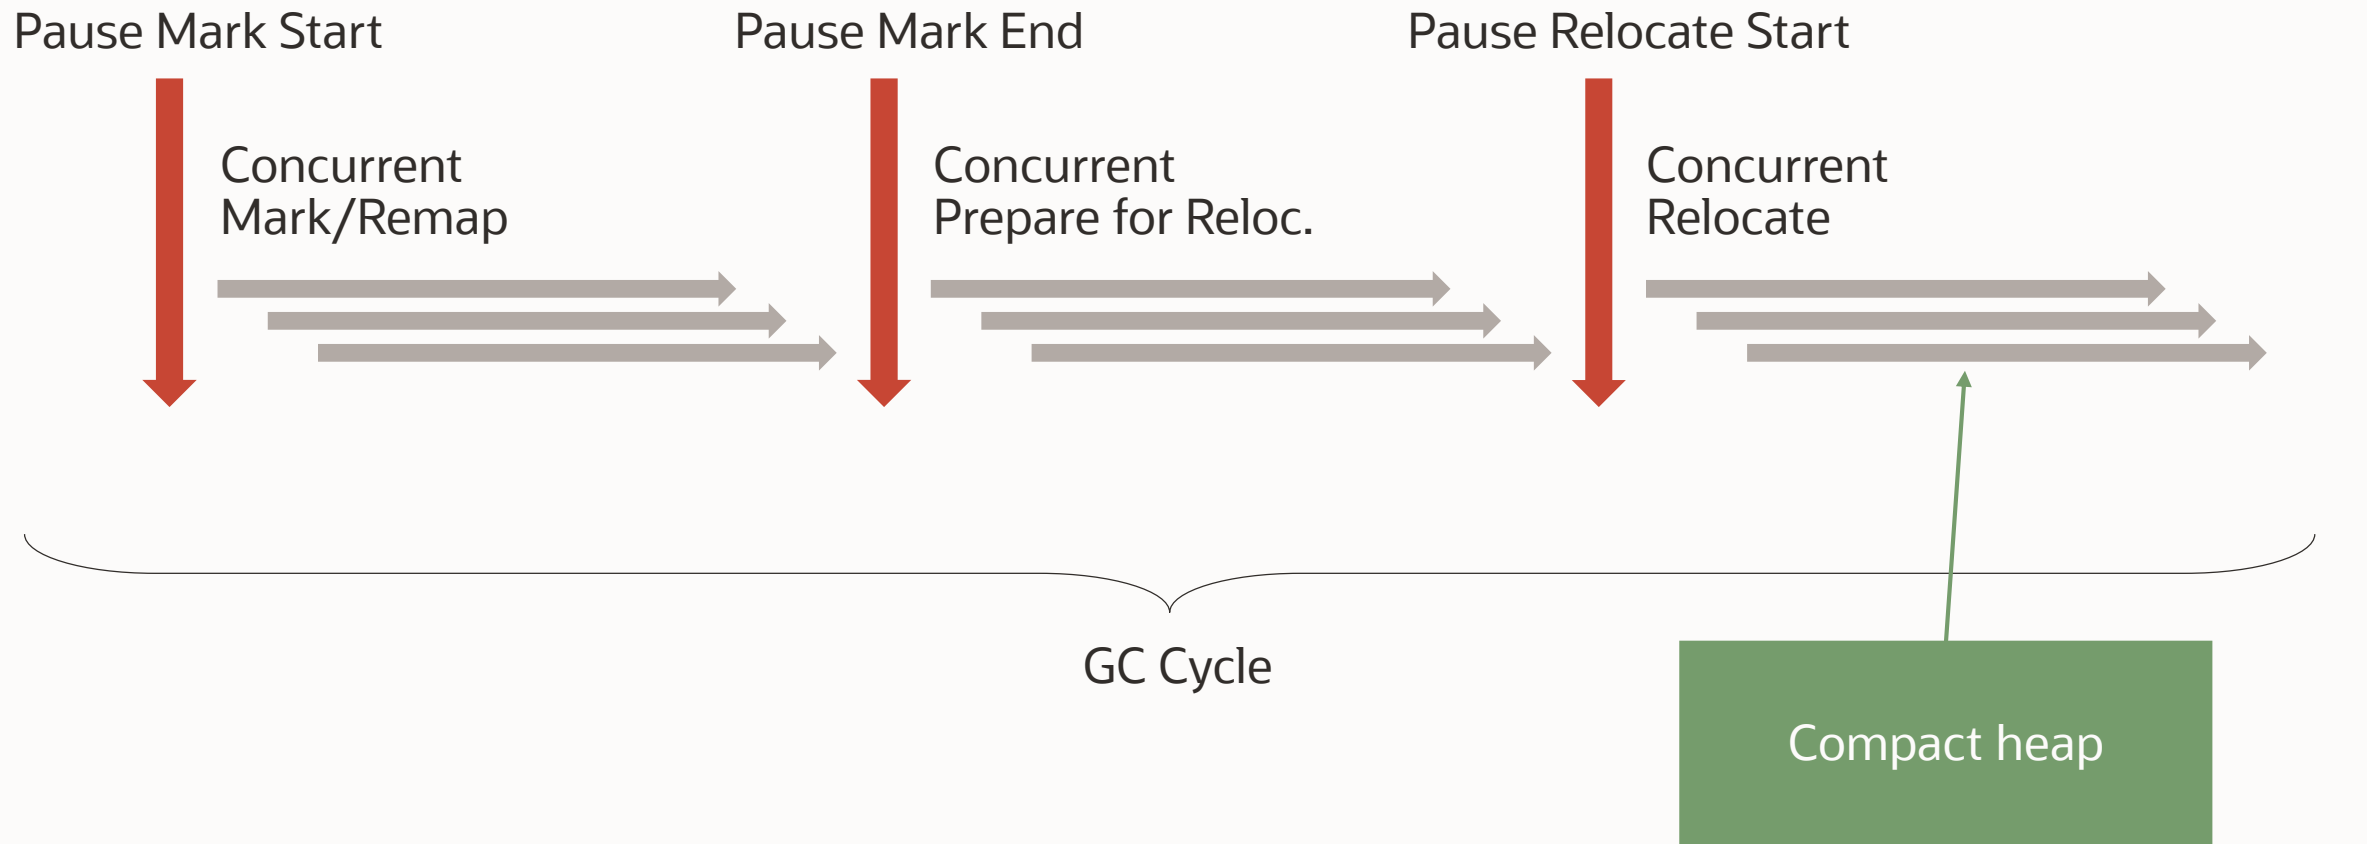


# ZGC Phases

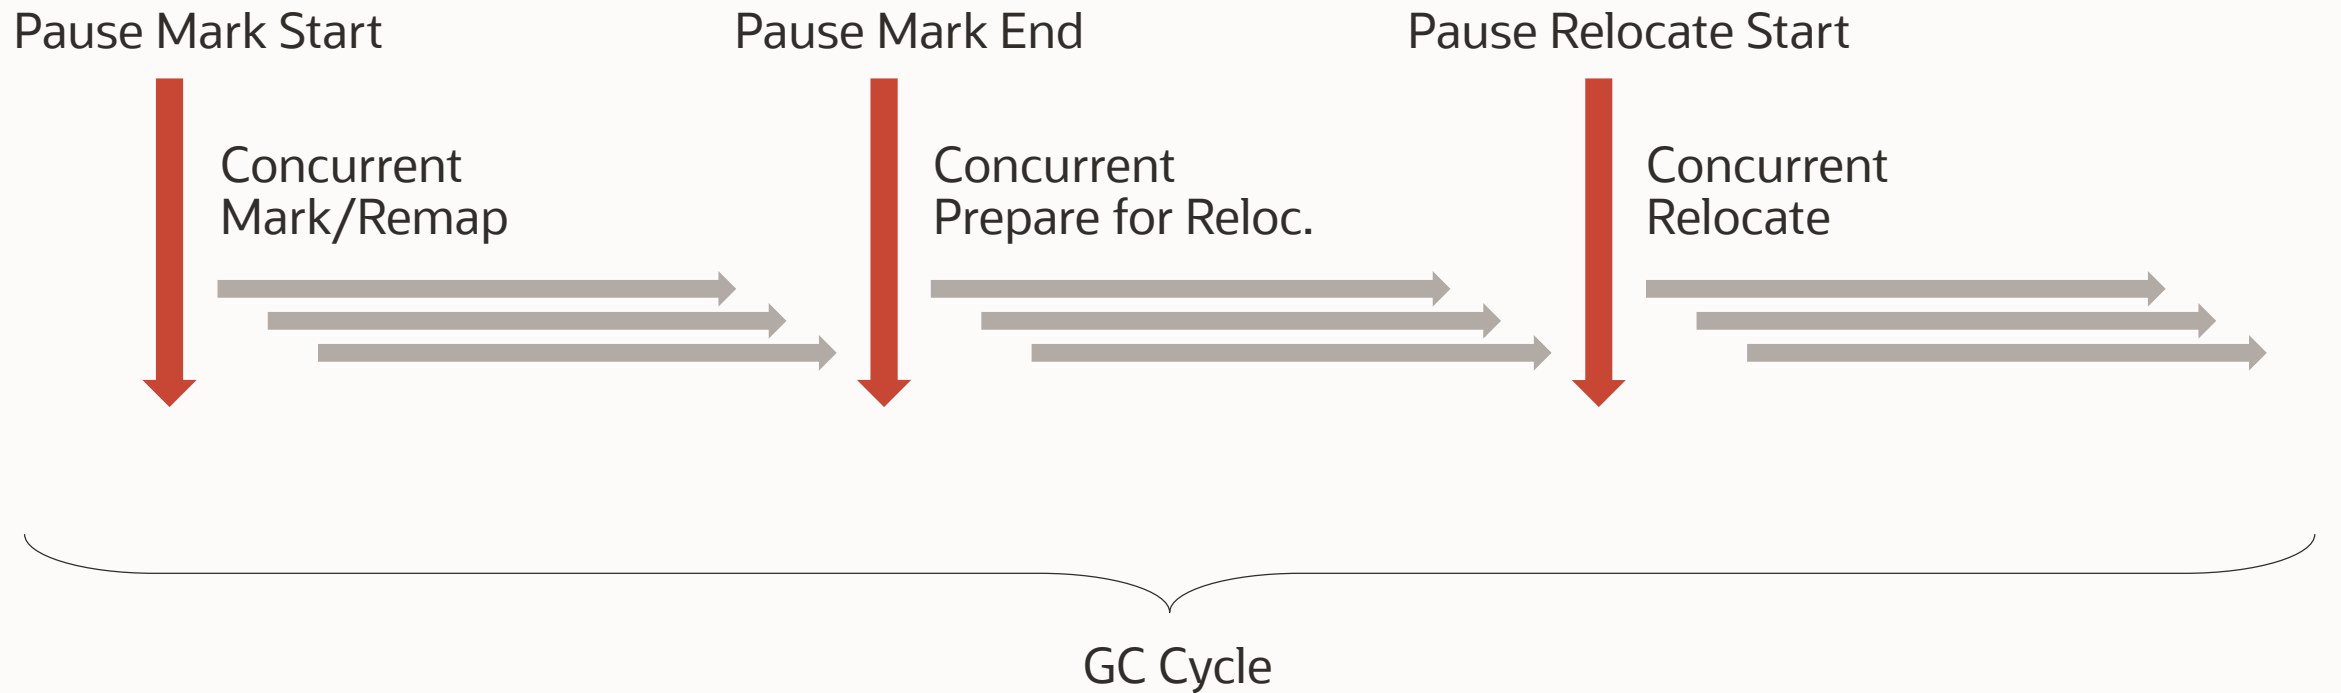




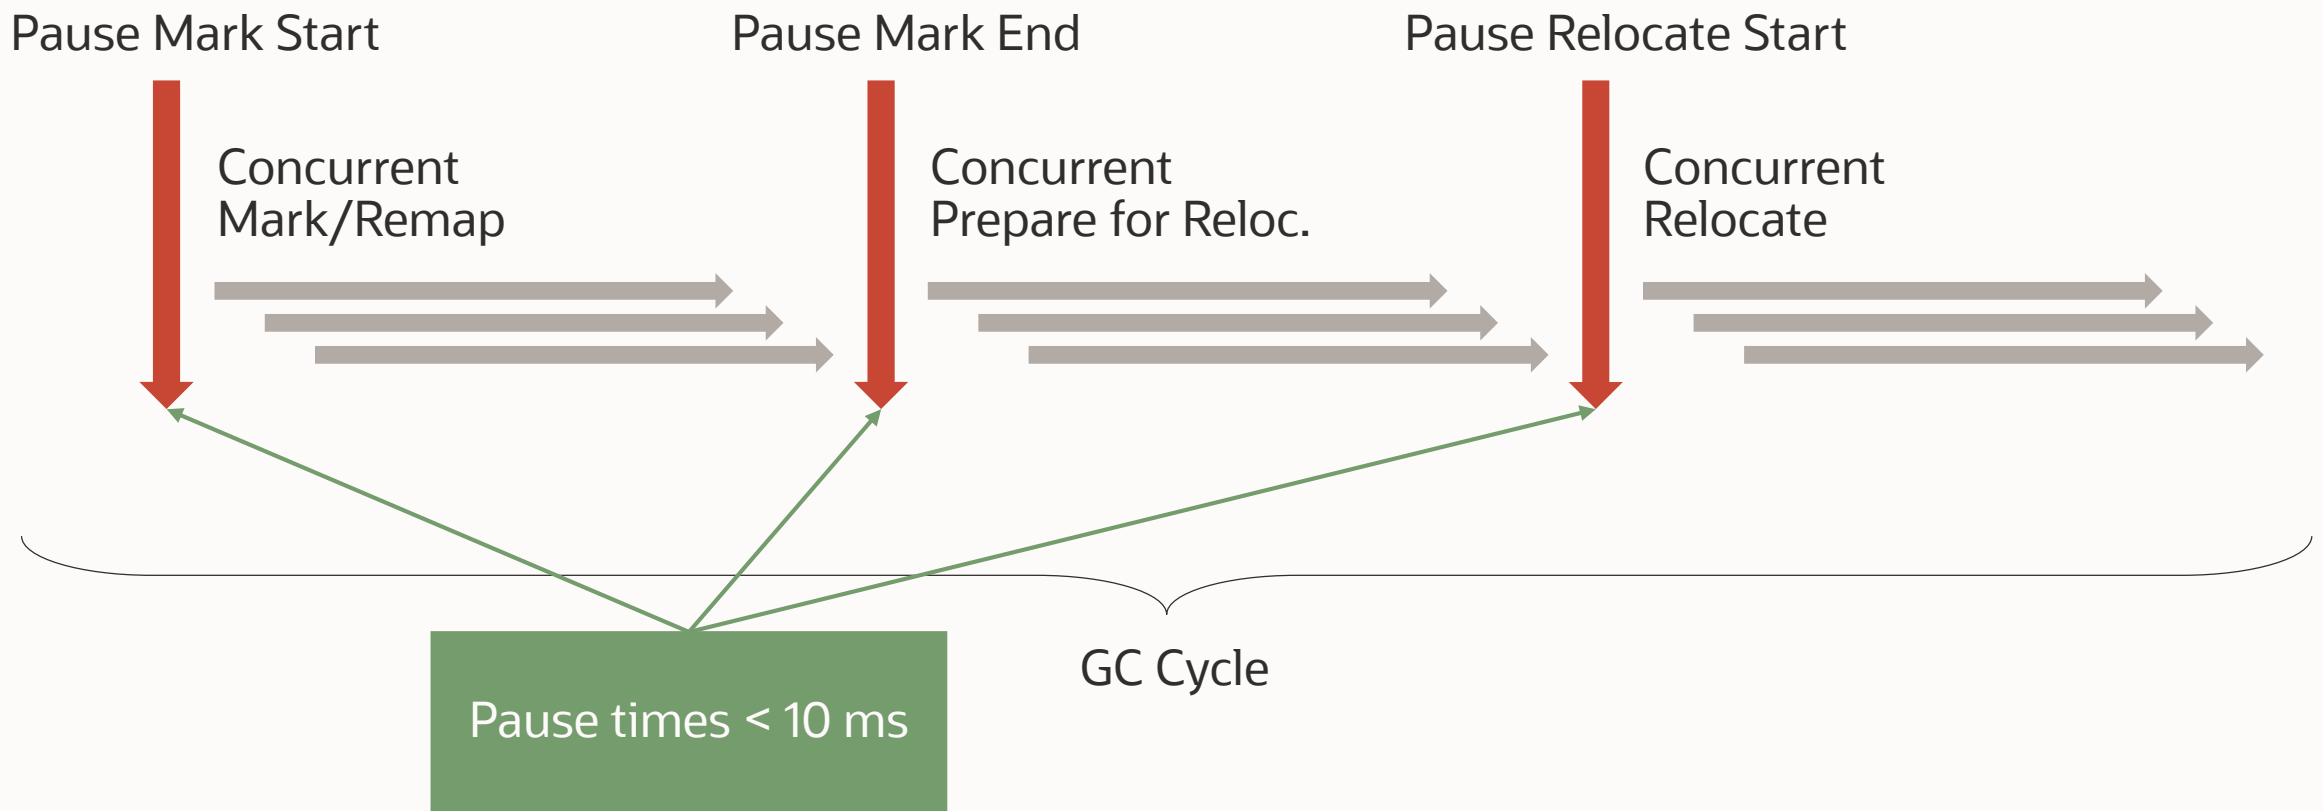
# ZGC Phases



# ZGC Phases

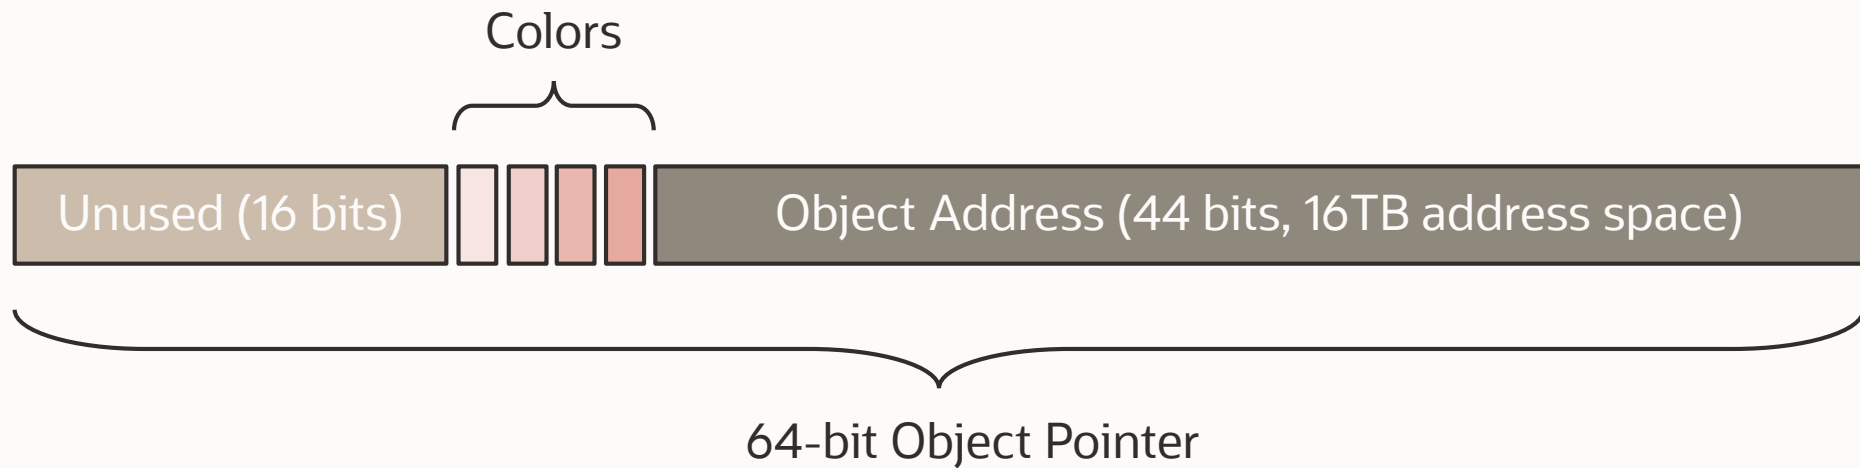


# ZGC Phases



# Colored Pointers

- Core design concept in ZGC
- **Metadata** stored in unused bits in 64-bit pointers



# Load Barrier



A small piece of code injected by the JIT in strategic places

- When **loading an object reference from the heap**

Checks if the loaded object reference has a **bad** color

- If so, take **action** and **heal** it

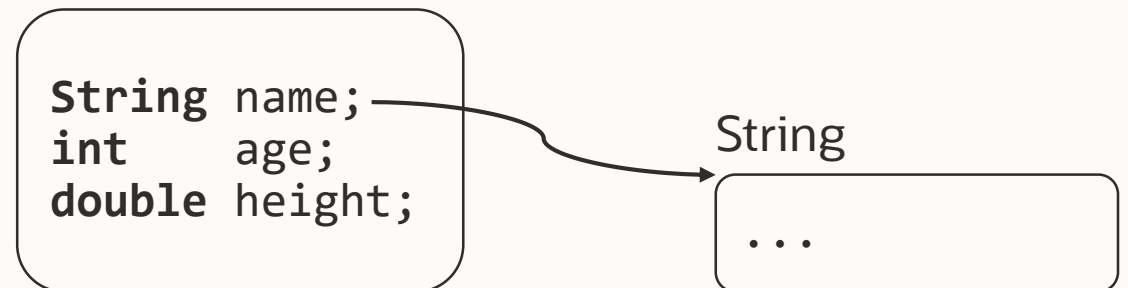
# Load Barrier



```
String n = person.name;
```

```
// Loading an object reference from heap
```

Person

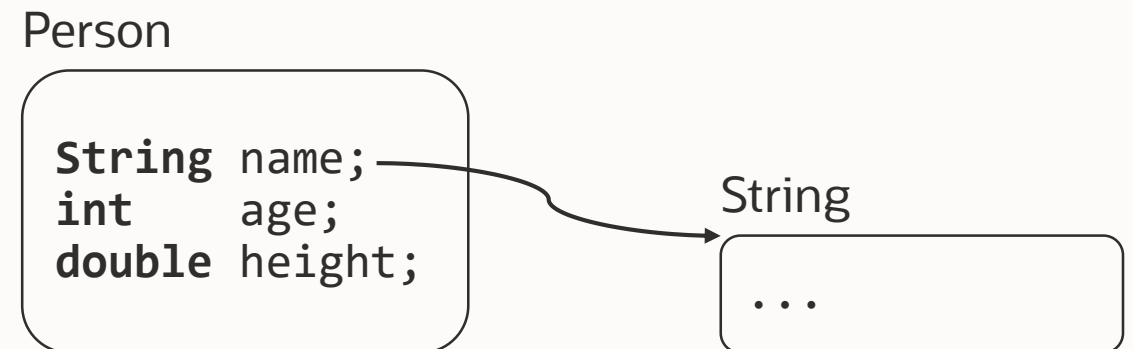


# Load Barrier



```
String n = person.name;  
<load barrier needed here>
```

```
// Loading an object reference from heap
```

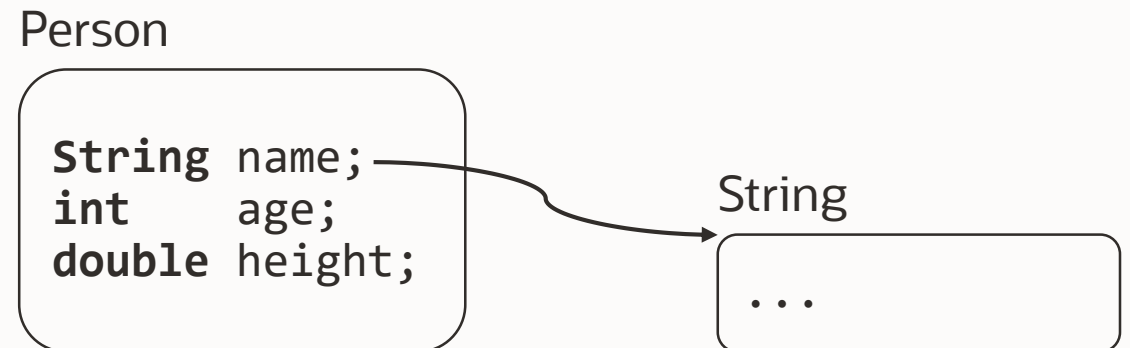


# Load Barrier



```
String n = person.name;  
<load barrier needed here>  
String p = n;  
n.isEmpty();  
int age = person.age;
```

```
// Loading an object reference from heap  
  
// No barrier, not a load from heap  
// No barrier, not a load from heap  
// No barrier, not an object reference
```



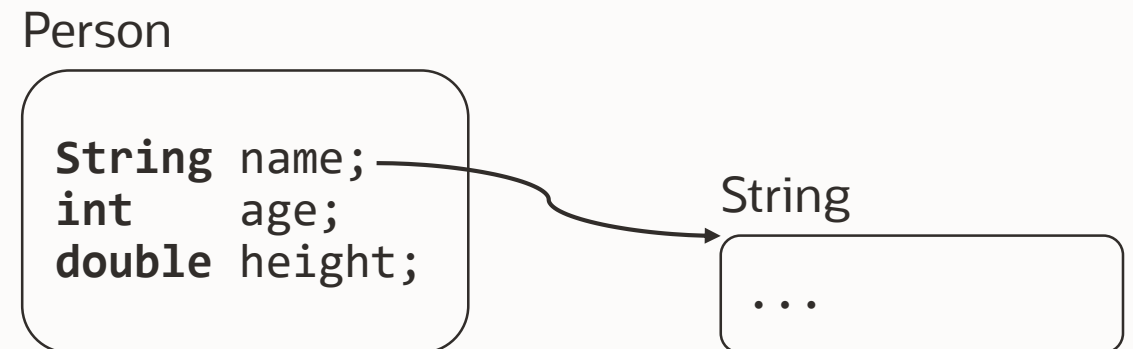


# Load Barrier



```
String n = person.name;  
<load barrier needed here>
```

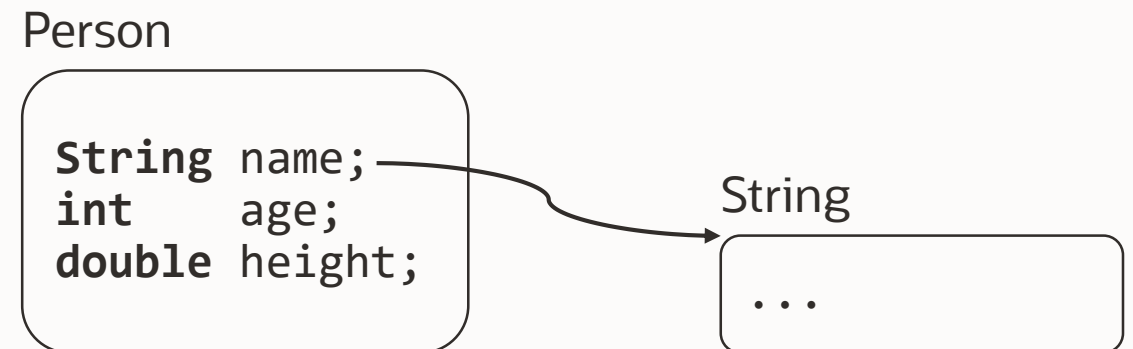
```
// Loading an object reference from heap
```



# Load Barrier



```
String n = person.name;           // Loading an object reference from heap
if (n & bad_bit_mask) {
    slow_path(register_for(n), address_of(person.name));
}
```



# Load Barrier



```
mov    0x10(%rax), %rbx
test   %rbx, 0x20(%r15)
jnz    slow_path
```

```
// String n = person.name;
// Bad color?
// Yes -> Enter slow path and
// mark/relocate/remap, adjust
// 0x10(%rax) and %rbx
```



# Performance

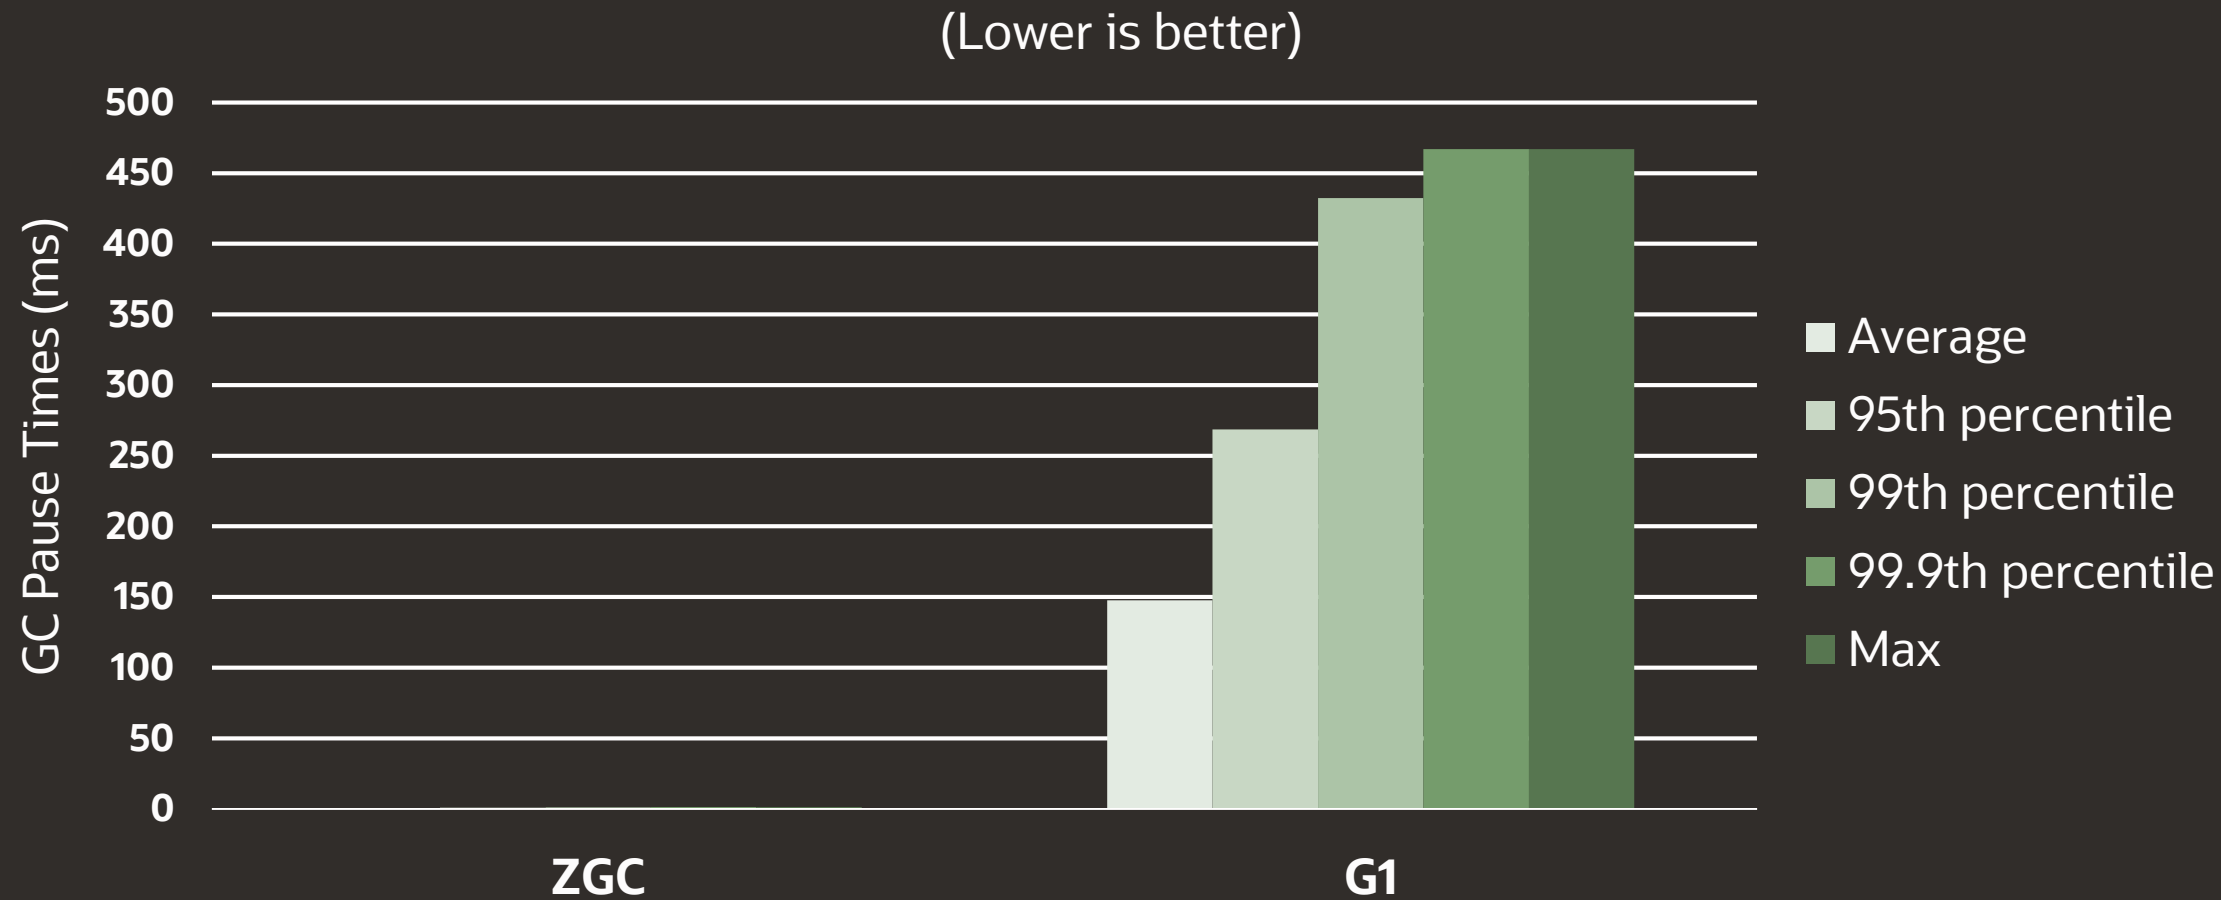
# SPECjbb2015 – Benchmark Score

128G Heap  
32 Hyper-threads (Intel)

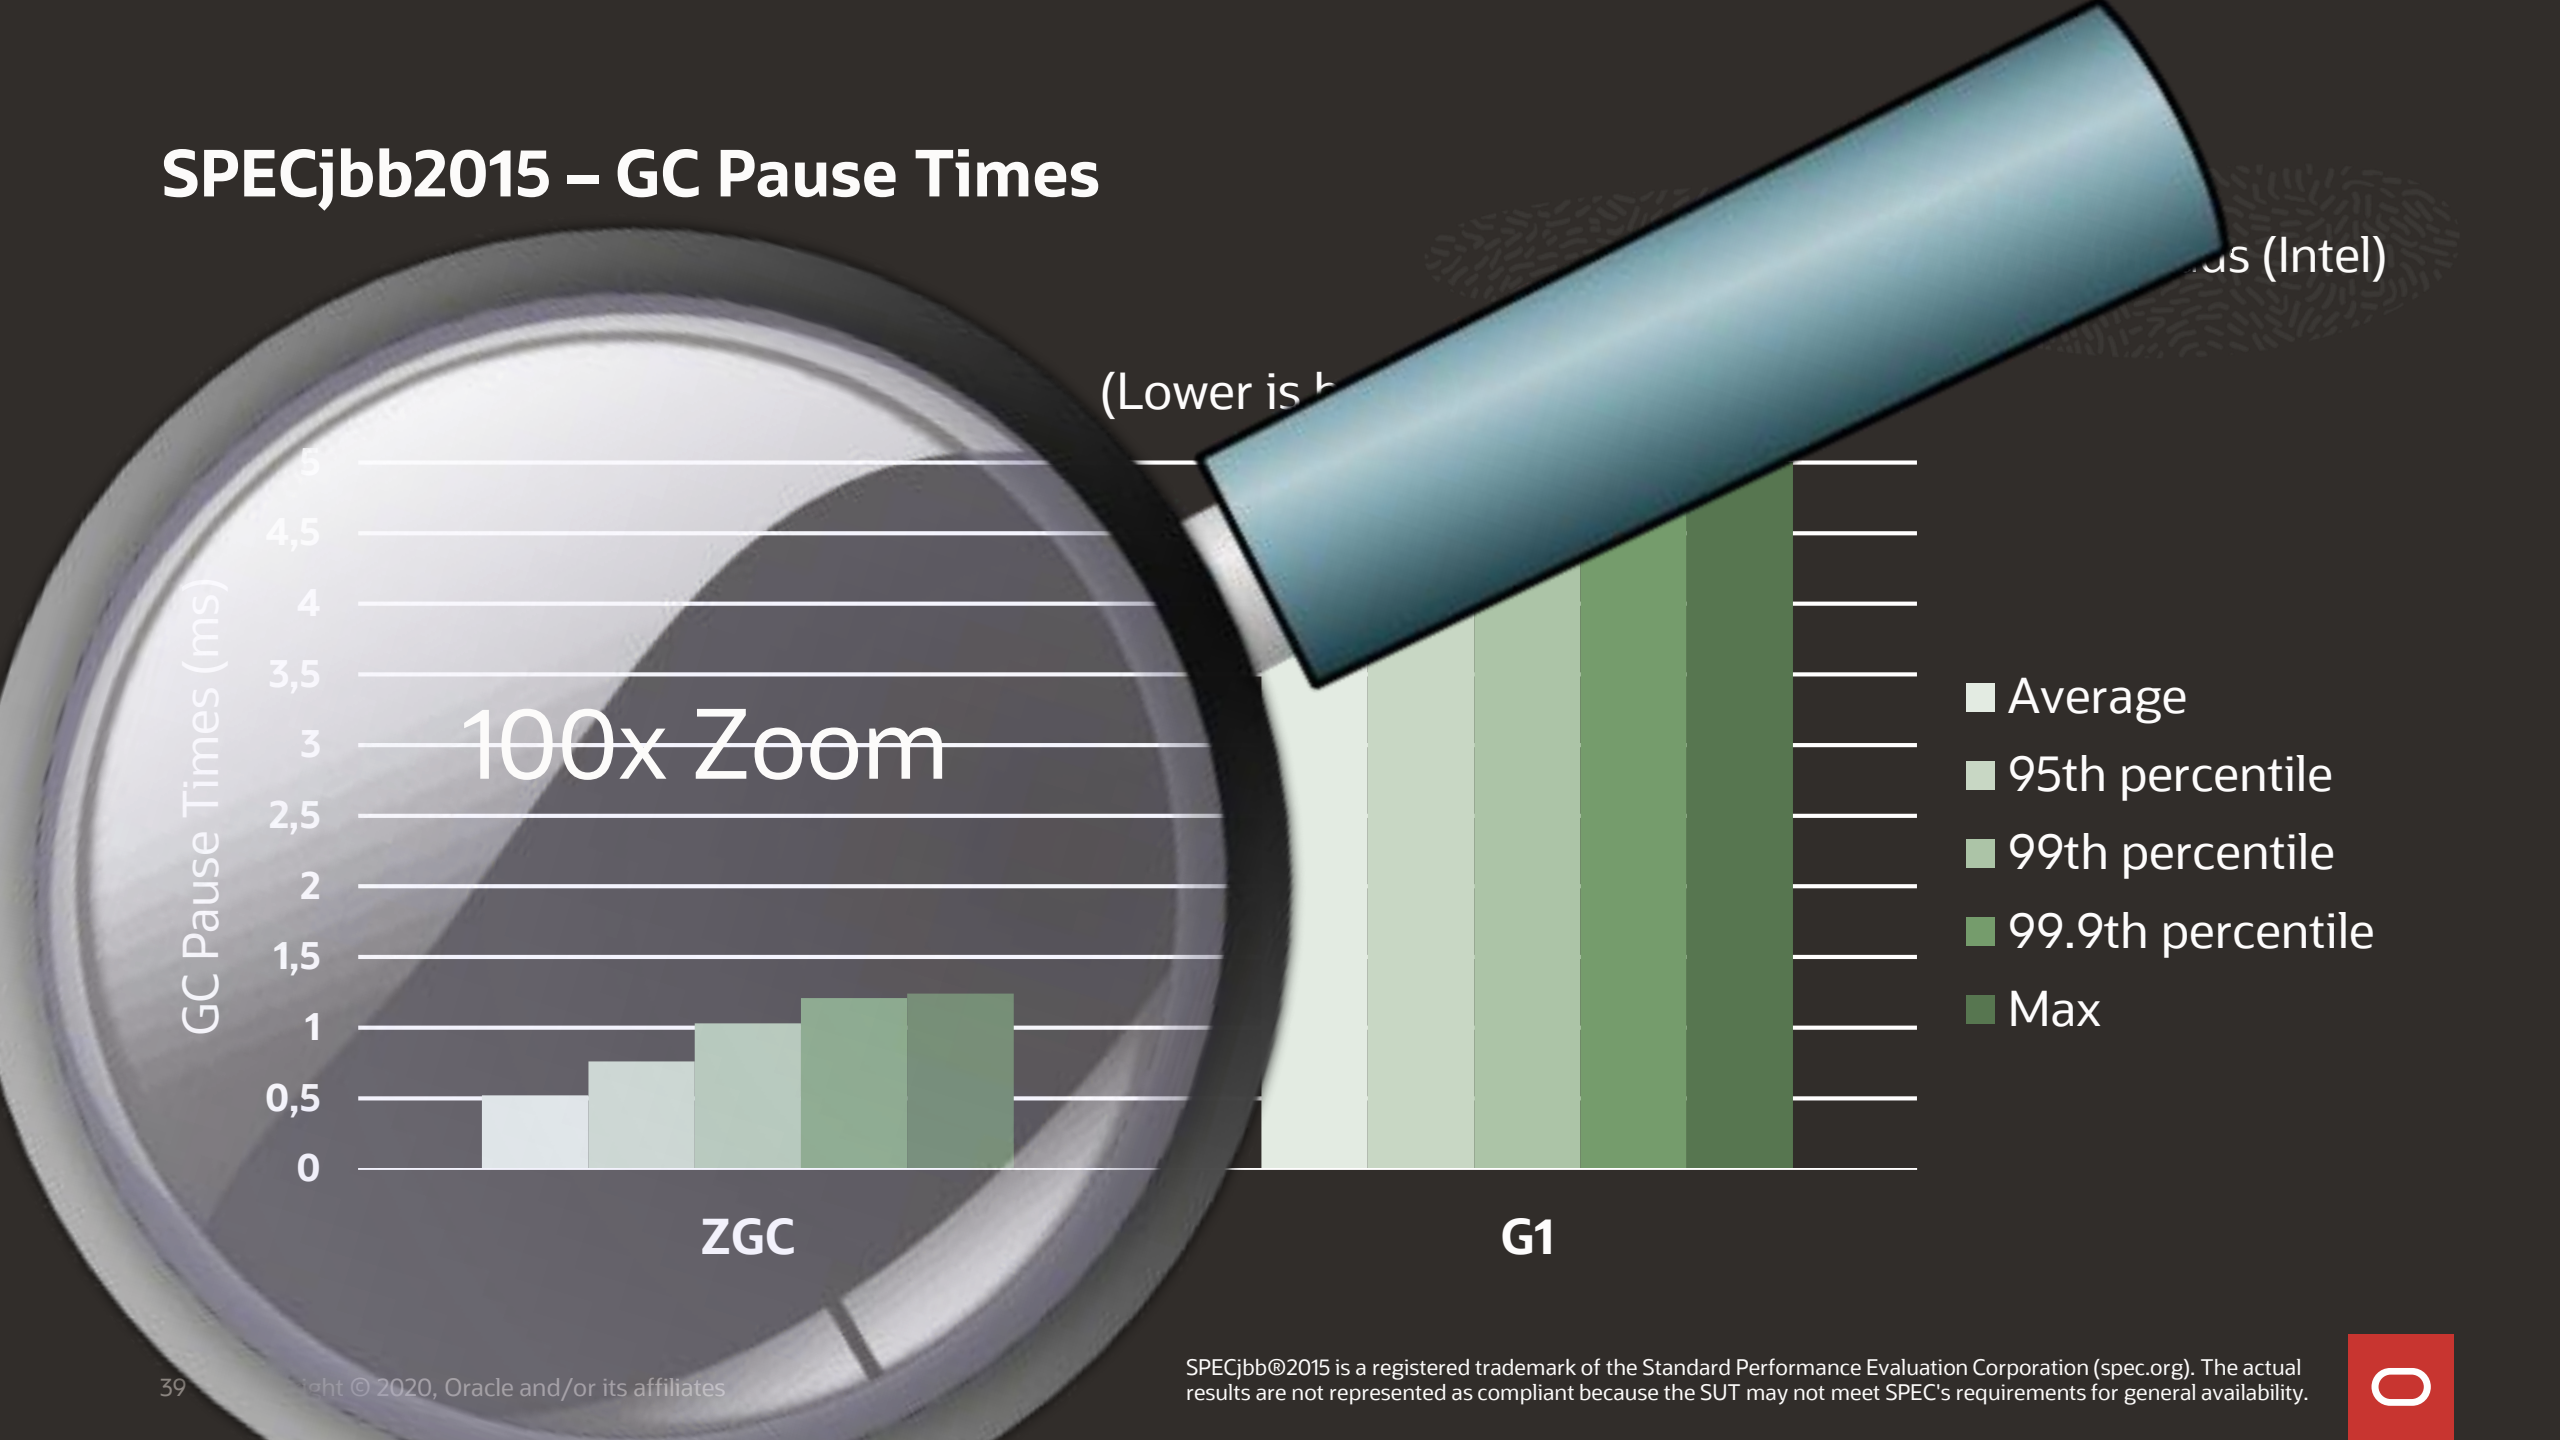


# SPECjbb2015 – GC Pause Times

128G Heap  
32 Hyper-threads (Intel)



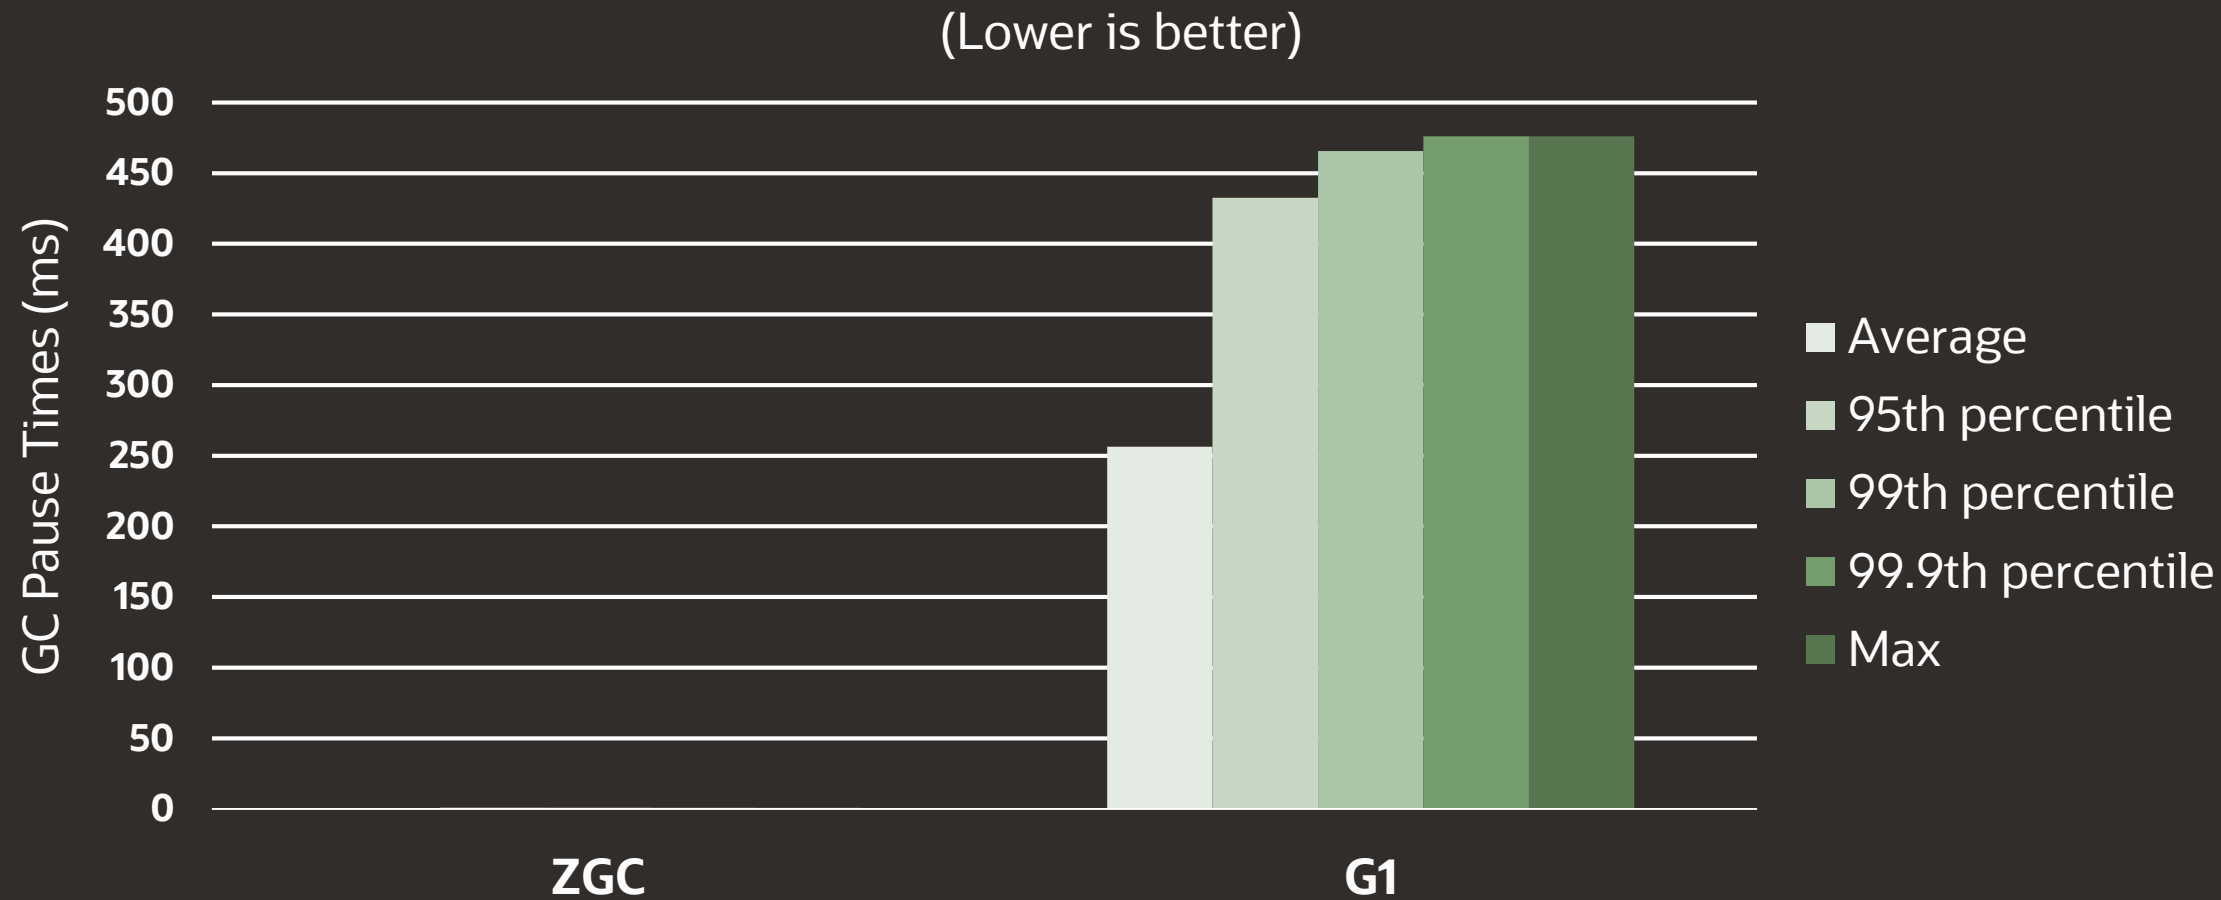
# SPECjbb2015 – GC Pause Times



# BigRamTester – GC Pause Times

Lots of heap fragmentation

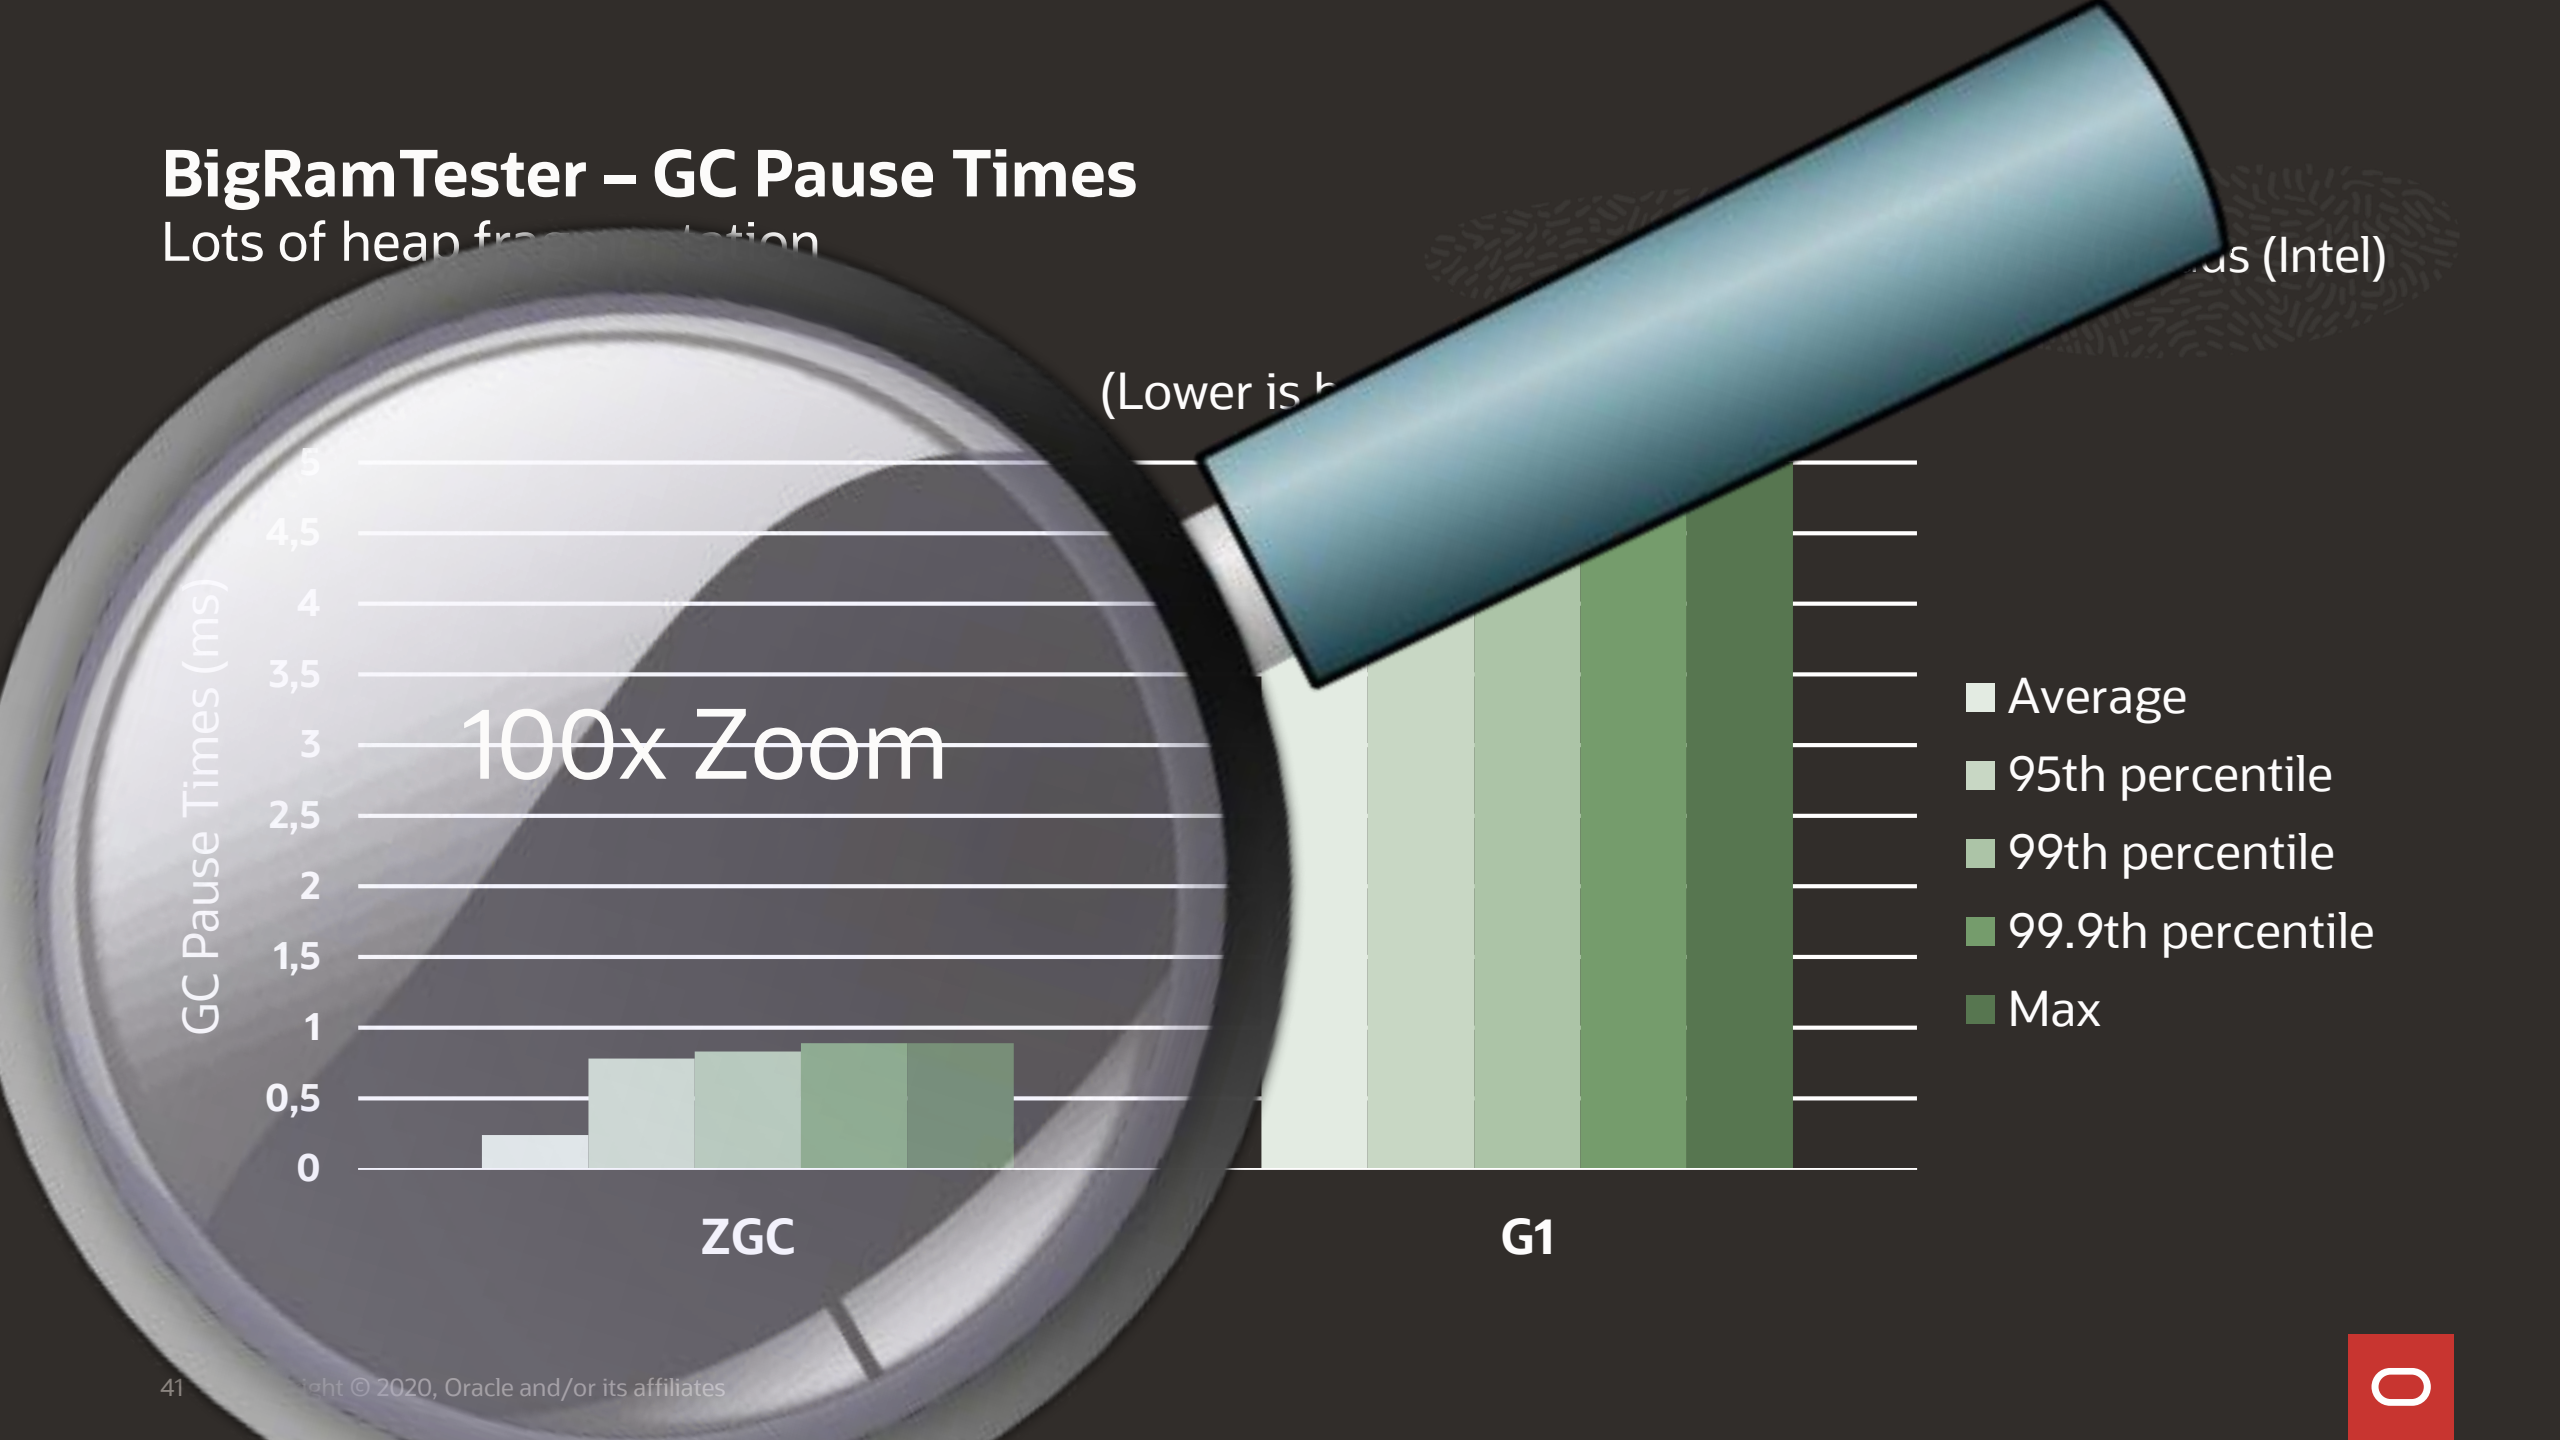
16G Heap  
32 Hyper-threads (Intel)





# BigRamTester – GC Pause Times

Lots of heap fragmentation





# Using ZGC

## Recommendation



Use the latest JDK version!

(Lots of ZGC improvements since **JDK 11**, and they will continue to come...)

# Enable



**-XX : +UseZGC**

# Enable



**-XX:+UnlockExperimentalVMOptions**  
**-XX:+UseZGC**

(JDK 14 and earlier)

# Tuning



# Tuning

## Set Max Heap Size



**-Xmx**<size>

# Tuning

*Maybe* Set Number of Concurrent GC Threads



`-Xmx<size>`

`-XX:ConcGCThreads=<number>`



# Tuning

That's it?

How much memory?

How much CPU-time?

**-Xmx**<size>

**-XX : ConcGCThreads=**<number>

# Logging



`-Xlog:gc` (basic)  
`-Xlog:gc*` (detailed)

Garbage Collection (Proactive) 13426M(10%) ->2492M(2%)  
 Garbage Collection (Allocation Rate) 87676M(67%) ->19578M(15%)  
 Garbage Collection (Allocation Rate) 55302M(42%) ->17646M(13%)  
 Garbage Collection (Allocation Rate) 61794M(47%) ->26794M(20%)  
 Garbage Collection (Allocation Rate) 60856M(46%) ->31926M(24%)  
 Garbage Collection (Allocation Rate) 52744M(40%) ->38050M(29%)  
 Garbage Collection (Allocation Rate) 42542M(32%) ->32204M(25%)  
 Garbage Collection (Allocation Rate) 49974M(38%) ->8534M(7%)  
 Garbage Collection (System.gc()) 8534M(7%) ->282M(0%)  
 Garbage Collection (Allocation Rate) 95454M(73%) ->25660M(20%)  
 Garbage Collection (Allocation Rate) 42478M(32%) ->23812M(18%)  
 Garbage Collection (Allocation Rate) 56714M(43%) ->29090M(22%)  
 Garbage Collection (Allocation Rate) 62802M(48%) ->28648M(22%)  
 Garbage Collection (Allocation Rate) 59748M(46%) ->23770M(18%)  
 Garbage Collection (Allocation Rate) 74946M(57%) ->23284M(18%)  
 Garbage Collection (System.gc()) 44902M(34%) ->422M(0%)  
 Garbage Collection (Allocation Rate) 94510M(72%) ->20456M(16%)  
 Garbage Collection (Allocation Rate) 59694M(46%) ->25834M(20%)  
 Garbage Collection (Allocation Rate) 63494M(48%) ->29128M(22%)  
 Garbage Collection (Allocation Rate) 59034M(45%) ->27094M(21%)  
 Garbage Collection (Allocation Rate) 66110M(50%) ->25278M(19%)  
 Garbage Collection (Allocation Rate) 73410M(56%) ->27968M(21%)  
 Garbage Collection (Allocation Rate) 70010M(53%) ->32236M(25%)  
 Garbage Collection (Allocation Rate) 64444M(49%) ->27612M(21%)  
 Garbage Collection (Allocation Rate) 64484M(49%) ->29910M(23%)  
 Garbage Collection (Allocation Rate) 64128M(49%) ->33184M(25%)  
 Garbage Collection (Allocation Rate) 59148M(45%) ->27800M(21%)  
 Garbage Collection (Allocation Rate) 63104M(48%) ->27976M(21%)  
 Garbage Collection (Allocation Rate) 64418M(49%) ->34390M(26%)  
 Garbage Collection (Allocation Rate) 52284M(40%) ->30654M(23%)  
 Garbage Collection (Allocation Rate) 58746M(45%) ->32028M(24%)  
 Garbage Collection (Allocation Rate) 59468M(45%) ->32804M(25%)  
 Garbage Collection (Allocation Rate) 53342M(41%) ->18436M(14%)



Garbage Collection (Proactive) 13426M(10%) ->2492M(2%)  
 Garbage Collection (Allocation Rate) 87676M(67%) ->19578M(15%)  
 Garbage Collection (Allocation Rate) 55302M(42%) ->17646M(13%)  
 Garbage Collection (Allocation Rate) 61794M(47%) ->26794M(20%)  
 Garbage Collection (Allocation Rate) 60856M(46%) ->31926M(24%)  
 Garbage Collection (Allocation Rate) 52744M(40%) ->38050M(29%)  
 Garbage Collection (Allocation Rate) 42542M(32%) ->32204M(25%)  
 Garbage Collection (Allocation Rate) 49974M(38%) ->8534M(7%)  
 Garbage Collection (System.gc()) 8534M(7%) ->282M(0%)  
**Garbage Collection (Allocation Rate) 95454M(73%) ->25660M(20%)**  
 Garbage Collection (Allocation Rate) 42478M(32%) ->23812M(18%)  
 Garbage Collection (Allocation Rate) 56714M(43%) ->29090M(22%)  
 Garbage Collection (Allocation Rate) 62802M(48%) ->28648M(22%)  
 Garbage Collection (Allocation Rate) 59748M(46%) ->23770M(18%)  
 Garbage Collection (Allocation Rate) 74946M(57%) ->23284M(18%)  
 Garbage Collection (System.gc()) 44902M(34%) ->422M(0%)  
 Garbage Collection (Allocation Rate) 94510M(72%) ->20456M(16%)  
 Garbage Collection (Allocation Rate) 59694M(46%) ->25834M(20%)  
 Garbage Collection (Allocation Rate) 63494M(48%) ->29128M(22%)  
 Garbage Collection (Allocation Rate) 59034M(45%) ->27094M(21%)  
 Garbage Collection  
 Garbage Collection  
 Garbage Collection  
 Garbage Collection  
 Garbage Collection (Allocation Rate) 64484M(49%) ->29910M(23%)  
 Garbage Collection (Allocation Rate) 64128M(49%) ->33184M(25%)  
 Garbage Collection (Allocation Rate) 59148M(45%) ->27800M(21%)  
 Garbage Collection (Allocation Rate) 63104M(48%) ->27976M(21%)  
 Garbage Collection (Allocation Rate) 64418M(49%) ->34390M(26%)  
 Garbage Collection (Allocation Rate) 52284M(40%) ->30654M(23%)  
 Garbage Collection (Allocation Rate) 58746M(45%) ->32028M(24%)  
 Garbage Collection (Allocation Rate) 59468M(45%) ->32804M(25%)  
 Garbage Collection (Allocation Rate) 53342M(41%) ->18436M(14%)

**Garbage Collection (Allocation Rate) 95454M(73%) ->25660M(20%)**

## Garbage Collection (Allocation Rate)

Pause Mark Start 0.483ms

Concurrent Mark 2885.207ms

Pause Mark End 0.439ms

Concurrent Process Non-Strong References 23.932ms

Concurrent Reset Relocation Set 5.682ms

Concurrent Select Relocation Set 215.098ms

Pause Relocate Start 0.531ms

Concurrent Relocate 1071.711ms

Load: 57.66/38.08/17.83

MMU: 2ms/41.5%, 5ms/76.6%, 10ms/87.3%, 20ms/92.1%, 50ms/96.1%, 100ms/98.0%

Mark: 4 stripe(s), 3 proactive flush(es), 1 terminate flush(es), 1 completion(s), 0 continuation(s)

NMethods: 5686 registered, 533 unregistered

Metaspace: 19M used, 20M capacity, 20M committed, 1042M reserved

Soft: 449128 encountered, 0 discovered, 0 enqueued

Weak: 4299 encountered, 2155 discovered, 119 enqueued

Final: 56 encountered, 6 discovered, 0 enqueued

Phantom: 89 encountered, 60 discovered, 0 enqueued

Small Pages: 26830 / 53660M(93%), Empty: 6768M(12%), Compacting: 46374M(80%)->700M(1%)

Medium Pages: 118 / 3776M(7%), Empty: 3520M(6%), Compacting: 256M(0%)->64M(0%)

Large Pages: 6 / 280M(0%), Empty: 0M(0%), Compacting: 0M(0%)->0M(0%)

Relocation: Successful

Min Capacity: 131072M(100%)

Max Capacity: 131072M(100%)

Soft Max Capacity: 131072M(100%)

	Mark Start	Mark End	Relocate Start	Relocate End	High	Low
Capacity:	131072M (100%)	131072M (100%)	131072M (100%)	131072M (100%)	131072M (100%)	131072M (100%)
Reserve:	72M (0%)	72M (0%)	72M (0%)	72M (0%)	72M (0%)	72M (0%)
Free:	73284M (56%)	53408M (41%)	61818M (47%)	100740M (77%)	101796M (78%)	53120M (41%)
Used:	57716M (44%)	77592M (59%)	69182M (53%)	30260M (23%)	77880M (59%)	29204M (22%)
Live:	-	1405M (1%)	1405M (1%)	1405M (1%)	-	-
Allocated:	-	19876M (15%)	21754M (17%)	29462M (22%)	-	-
Garbage:	-	56310M (43%)	46022M (35%)	4990M (4%)	-	-
Reclaimed:	-	-	10288M (8%)	51320M (39%)	-	-

Garbage Collection (Allocation Rate) 57716M(44%)->30260M(23%)



## Garbage Collection (Allocation Rate)

Pause Mark Start 0.483ms  
Concurrent Mark 2885.207ms  
Pause Mark End 0.439ms  
Concurrent Process Non-Strong References 23.932ms  
Concurrent Reset Relocation Set 5.682ms  
Concurrent Select Relocation Set 215.098ms  
Pause Relocate Start 0.531ms  
Concurrent Relocate 1071.711ms

Load: 57.66/38.08/17.83  
MMU: 2ms/41.5%, 5ms/76.6%, 10ms/87.3%, 20ms/93.8%  
Mark: 4 stripe(s), 3 proactive flush(es), 1  
NMethods: 5686 registered, 533 unregistered  
Metaspace: 19M used, 20M capacity, 20M committed  
Soft: 449128 encountered, 0 discovered, 0 enqueued  
Weak: 4299 encountered, 2155 discovered, 119 enqueued  
Final: 56 encountered, 6 discovered, 0 enqueued  
Phantom: 89 encountered, 60 discovered, 0 enqueued  
Small Pages: 26830 / 53660M(93%), Empty: 67680M(6%)  
Medium Pages: 118 / 3776M(7%), Empty: 3520M(9%)  
Large Pages: 6 / 280M(0%), Empty: 0M(0%), Committed: 168M(6%)  
Relocation: Successful  
Min Capacity: 131072M(100%)  
Max Capacity: 131072M(100%)  
Soft Max Capacity: 131072M(100%)

Pause Mark Start 0.483ms  
Concurrent Mark 2885.207ms  
Pause Mark End 0.439ms  
Concurrent Process Non-Strong References 23.932ms  
Concurrent Reset Relocation Set 5.682ms  
Concurrent Select Relocation Set 215.098ms  
Pause Relocate Start 0.531ms  
Concurrent Relocate 1071.711ms

	Mark Start	Mark End	Relocate Start	Relocate End	High	Low
Capacity:	131072M (100%)	131072M (100%)	131072M (100%)	131072M (100%)	131072M (100%)	131072M (100%)
Reserve:	72M (0%)	72M (0%)	72M (0%)	72M (0%)	72M (0%)	72M (0%)
Free:	73284M (56%)	53408M (41%)	61818M (47%)	100740M (77%)	101796M (78%)	53120M (41%)
Used:	57716M (44%)	77592M (59%)	69182M (53%)	30260M (23%)	77880M (59%)	29204M (22%)
Live:	-	1405M (1%)	1405M (1%)	1405M (1%)	-	-
Allocated:	-	19876M (15%)	21754M (17%)	29462M (22%)	-	-
Garbage:	-	56310M (43%)	46022M (35%)	4990M (4%)	-	-
Reclaimed:	-	-	10288M (8%)	51320M (39%)	-	-
Garbage Collection (Allocation Rate) 57716M(44%)->30260M(23%)						

Garbage Collection (Allocation Rate)  
Pause Mark Start 0.483ms  
Concurrent Mark 2885.207ms  
Pause Mark End 0.439ms  
Concurrent Process Non-Strong References 23.932ms  
Concurrent Reset Relocation Set 5.682ms  
Concurrent Select  
Pause Relocate Start  
Concurrent Relocate  
Load: 57.66/38.08  
MMU: 2ms/41.5%,  
Mark: 4 stripe(s)  
NMethods: 5686 re  
Metaspace: 19M us  
Soft: 449128 encou  
Weak: 4299 encour  
Final: 56 encount  
Phantom: 89 encou  
Small Pages: 268  
Medium Pages: 118  
Large Pages: 6 / 280M(0%), Empty: 0M(0%), Compacting: 0M(0%)->0M(0%)  
Relocation: Successful  
Min Capacity: 131072M(100%)  
Max Capacity: 131072M(100%)  
Soft Max Capacity: 131072M(100%)

	Mark Start	Mark End	Relocate Start	Relocate End
Capacity:	131072M (100%)	131072M (100%)	131072M (100%)	131072M (100%)
Reserve:	72M (0%)	72M (0%)	72M (0%)	72M (0%)
Free:	73284M (56%)	53408M (41%)	61818M (47%)	100740M (77%)
Used:	57716M (44%)	77592M (59%)	69182M (53%)	30260M (23%)
Live:	-	1405M (1%)	1405M (1%)	1405M (1%)
Allocated:	-	19876M (15%)	21754M (17%)	29462M (22%)
Garbage:	-	56310M (43%)	46022M (35%)	4990M (4%)
Reclaimed:	-	-	10288M (8%)	51320M (39%)

	Mark Start	Mark End	Relocate Start	Relocate End
Capacity:	131072M (100%)	131072M (100%)	131072M (100%)	131072M (100%)
Reserve:	72M (0%)	72M (0%)	72M (0%)	72M (0%)
Free:	73284M (56%)	53408M (41%)	61818M (47%)	100740M (77%)
Used:	57716M (44%)	77592M (59%)	69182M (53%)	30260M (23%)
Live:	-	1405M (1%)	1405M (1%)	1405M (1%)
Allocated:	-	19876M (15%)	21754M (17%)	29462M (22%)
Garbage:	-	56310M (43%)	46022M (35%)	4990M (4%)
Reclaimed:	-	-	10288M (8%)	51320M (39%)

High	Low
131072M (100%)	131072M (100%)
72M (0%)	72M (0%)
101796M (78%)	53120M (41%)
77880M (59%)	29204M (22%)
-	-
-	-
-	-
-	-

Garbage Collection (Allocation Rate) 57716M(44%)->30260M(23%)



## Garbage Collection (Allocation Rate)

Pause Mark Start 0.483ms

Concurrent Mark 2885.207ms

Pause Mark End 0.439ms

Concurrent Process Non-Strong References 23.932ms

Concurrent Reset Relocation Set 5.682ms

Concurrent Select Relocation Set 215.098ms

Pause Relocate Start 0.531ms

Concurrent Relocate 1071.711ms

Load: 57.66/38.08/17.83

MMU: 2ms/41.5%, 5ms/76.6%, 10ms/87.3%, 20ms/92.1%, 50ms/96.1%, 100ms/98.0%

Mark: 4 stripe(s), 3 proactive flush(es), 1 terminate flush(es), 1 completion(s), 0 continuation(s)

NMethods: 5686 registered, 533 unregistered

Metaspace: 19M used, 20M capacity, 20M committed, 1042M reserved

Soft: 449128 encountered, 0 discovered, 0 enqueued

Weak: 4299 encountered, 2155 discovered, 119 enqueued

Final: 56 encountered, 6 discovered, 0 enqueued

Phantom: 89 encountered, 60 discovered, 0 enqueued

Small Pages: 26830 / 53660M(93%), Empty: 6768M(12%), Compacting: 46374M(80%)->700M(1%)

Medium Pages: 118 / 3776M(7%), Empty: 3520M(6%), Compacting: 256M(0%)->64M(0%)

Large Pages: 6 / 280M(0%), Empty: 0M(0%), Compacting: 0M(0%)->0M(0%)

Relocation: Successful

Min Capacity: 131072M(100%)

Max Capacity: 131072M(100%)

Soft Max Capacity: 131072M(100%)

	Mark Start	Mark End	Relocate Start	Relocate End	High	Low
Capacity:	131072M (100%)	131072M (100%)	131072M (100%)	131072M (100%)	131072M (100%)	131072M (100%)
Reserve:	72M (0%)	72M (0%)	72M (0%)	72M (0%)	72M (0%)	72M (0%)
Free:	73284M (56%)	53408M (41%)	61818M (47%)	100740M (77%)	101796M (78%)	53120M (41%)
Used:	57716M (44%)	77592M (59%)	69182M (53%)	30260M (23%)	77880M (59%)	29204M (22%)
Live:	-	1405M (1%)	1405M (1%)	1405M (1%)	-	-
Allocated:	-	19876M (15%)	21754M (17%)	29462M (22%)	-	-
Garbage:	-	56310M (43%)	46022M (35%)	4990M (4%)	-	-
Reclaimed:	-	-	10288M (8%)	51320M (39%)	-	-

Garbage Collection (Allocation Rate) 57716M(44%)->30260M(23%)





=== Garbage Collection Statistics =====					
	Last 10s	Last 10m	Last 10h	Total	
	Avg / Max	Avg / Max	Avg / Max	Avg / Max	
Collector: Garbage Collection Cycle	393.373 / 393.373	2457.026 / 4420.944	2457.026 / 4420.944	2457.026 / 4420.944	ms
Contention: Mark Segment Reset Contention	4 / 47	1 / 165	1 / 165	1 / 165	ops/s
Contention: Mark SeqNum Reset Contention	0 / 0	0 / 1	0 / 1	0 / 1	ops/s
Contention: Relocation Contention	0 / 6	13 / 1609	13 / 1609	13 / 1609	ops/s
Critical: Allocation Stall	0.000 / 0.000	0.000 / 0.000	0.000 / 0.000	0.000 / 0.000	ms
Critical: Allocation Stall	0 / 0	0 / 0	0 / 0	0 / 0	ops/s
Critical: GC Locker Stall	0.000 / 0.000	0.000 / 0.000	0.000 / 0.000	0.000 / 0.000	ms
Critical: GC Locker Stall	0 / 0	0 / 0	0 / 0	0 / 0	ops/s
Memory: Allocation Rate	862 / 2892	2229 / 13172	2229 / 13172	2229 / 13172	MB/s
Memory: Heap Used After Mark	6038 / 6038	55657 / 107472	55657 / 107472	55657 / 107472	MB
Memory: Heap Used After Relocation	1496 / 1496	14059 / 31084	14059 / 31084	14059 / 31084	MB
Memory: Heap Used Before Mark	6038 / 6038	47444 / 94452	47444 / 94452	47444 / 94452	MB
Memory: Heap Used Before Relocation	3452 / 3452	47152 / 94430	47152 / 94430	47152 / 94430	MB
Memory: Out Of Memory	0 / 0	0 / 0	0 / 0	0 / 0	ops/s
Memory: Page Cache Flush	0 / 0	0 / 0	0 / 0	0 / 0	MB/s
Memory: Page Cache Hit L1	350 / 1183	765 / 3549	765 / 3549	765 / 3549	ops/s
Memory: Page Cache Hit L2	0 / 0	4 / 234	4 / 234	4 / 234	ops/s
Memory: Page Cache Hit L3	2 / 7	127 / 2928	127 / 2928	127 / 2928	ops/s
Memory: Page Cache Miss	0 / 0	0 / 0	0 / 0	0 / 0	ops/s
Memory: Uncommit	0 / 0	0 / 0	0 / 0	0 / 0	MB/s
Memory: Undo Object Allocation Failed	0 / 0	0 / 53	0 / 53	0 / 53	ops/s
Memory: Undo Object Allocation Succeeded	0 / 6	12 / 1609	12 / 1609	12 / 1609	ops/s
Memory: Undo Page Allocation	0 / 0	6 / 463	6 / 463	6 / 463	ops/s
Phase: Concurrent Mark	317.841 / 317.841	1564.321 / 3056.192	1564.321 / 3056.192	1564.321 / 3056.192	ms
Phase: Concurrent Mark Continue	0.000 / 0.000	8.173 / 8.173	8.173 / 8.173	8.173 / 8.173	ms
Phase: Concurrent Process Non-Strong References	1.533 / 1.533	16.783 / 57.370	16.783 / 57.370	1783 / 57.370	ms
Phase: Concurrent Relocate	49.909 / 49.909	748.800 / 1299.535	748.800 / 1299.535	748.800 / 1299.535	ms
Phase: Concurrent Reset Relocation Set	5.138 / 5.138	4.566 / 12.218	4.566 / 12.218	4.566 / 12.218	ms
Phase: Concurrent Select Relocation Set	16.672 / 16.672	111.356 / 259.445	111.356 / 259.445	111.356 / 259.445	ms
Phase: Pause Mark End	0.063 / 0.063	0.411 / 1.142	0.411 / 1.142	0.411 / 1.142	ms
Phase: Pause Mark Start	0.250 / 0.250	0.407 / 0.706	0.407 / 0.706	0.407 / 0.706	ms
Phase: Pause Relocate Start	0.513 / 0.513	0.598 / 1.170	0.598 / 1.170	0.598 / 1.170	ms
Subphase: Concurrent Classes Purge	0.112 / 0.112	0.293 / 0.709	0.293 / 0.709	0.293 / 0.709	ms
Subphase: Concurrent Classes Unlink	0.897 / 0.897	5.852 / 18.071	5.852 / 18.071	5.852 / 18.071	ms
Subphase: Concurrent Mark	315.060 / 317.380	814.912 / 3055.673	814.912 / 3055.673	814.912 / 3055.673	ms
Subphase: Concurrent Mark Idle	1.075 / 1.161	1.958 / 21.000	1.958 / 21.000	1.958 / 21.000	ms
Subphase: Concurrent Mark Try Flush	0.137 / 0.449	6.628 / 51.534	6.628 / 51.534	6.628 / 51.534	ms
...					



=== Garbage Collection Statistics =====

	Last 10s	Last 10m	Last 10h	Total	
	Avg / Max	Avg / Max	Avg / Max	Avg / Max	
Collector: Garbage Collection Cycle	393.373 / 393.373	2457.026 / 4420.944	2457.026 / 4420.944	2457.026 / 4420.944	ms
Contention: Mark Segment Reset Contention	4 / 47	1 / 165	1 / 165	1 / 165	ops/s
Contention: Mark SeqNum Reset Contention	0 / 0	0 / 1	0 / 1	0 / 1	ops/s
Contention: Relocation Contention	0 / 6	13 / 1609	13 / 1609	13 / 1609	ops/s
Critical: Allocation Stall	0.000 / 0.000	0.000 / 0.000	0.000 / 0.000	0.000 / 0.000	ms
Critical: Allocation Stall	0 / 0	0 / 0	0 / 0	0 / 0	ops/s
Critical: GC Locker Stall	0.000 / 0.000	0.000 / 0.000	0.000 / 0.000	0.000 / 0.000	ms
Critical: GC Locker Stall	0 / 0	0 / 0	0 / 0	0 / 0	ops/s
Memory: Allocation Rate	862 / 2892	2229 / 13172	2229 / 13172	2229 / 13172	MB/s
Memory: Heap Used After Mark	6038 / 6038	55657 / 107472	55657 / 107472	55657 / 107472	MB
Memory: Heap Used After Relocation	1496 / 1496	14059 / 31084	14059 / 31084	14059 / 31084	MB
Memory: Heap Used Before Mark	6038 / 6038	47444 / 94452	47444 / 94452	47444 / 94452	MB
Memory: Heap Used Before Relocation	3452 / 3452	47152 / 94430	47152 / 94430	47152 / 94430	MB
Memory: Out Of Memory	0 / 0	0 / 0	0 / 0	0 / 0	ops/s
Memory: Page Cache Flush	0 / 0	0 / 0	0 / 0	0 / 0	MB/s
Memory: Page Cache Hit L1	350 / 1183	765 / 3549	765 / 3549	765 / 3549	ops/s

	Last 10s	Last 10m	Last 10h	Total	
	Avg / Max	Avg / Max	Avg / Max	Avg / Max	
Allocation Rate	862 / 2892	2229 / 13172	2229 / 13172	2229 / 13172	MB/s
Memory: Undo Page Allocation	0 / 0	6 / 463	6 / 463	6 / 463	ops/s
Phase: Concurrent Mark	317.841 / 317.841	1564.321 / 3056.192	1564.321 / 3056.192	1564.321 / 3056.192	ms
Phase: Concurrent Mark Continue	0.000 / 0.000	8.173 / 8.173	8.173 / 8.173	8.173 / 8.173	ms
Phase: Concurrent Process Non-Strong References	1.533 / 1.533	16.783 / 57.370	16.783 / 57.370	1783 / 57.370	ms
Phase: Concurrent Relocate	49.909 / 49.909	748.800 / 1299.535	748.800 / 1299.535	748.800 / 1299.535	ms
Phase: Concurrent Reset Relocation Set	5.138 / 5.138	4.566 / 12.218	4.566 / 12.218	4.566 / 12.218	ms
Phase: Concurrent Select Relocation Set	16.672 / 16.672	111.356 / 259.445	111.356 / 259.445	111.356 / 259.445	ms
Phase: Pause Mark End	0.063 / 0.063	0.411 / 1.142	0.411 / 1.142	0.411 / 1.142	ms
Phase: Pause Mark Start	0.250 / 0.250	0.407 / 0.706	0.407 / 0.706	0.407 / 0.706	ms
Phase: Pause Relocate Start	0.513 / 0.513	0.598 / 1.170	0.598 / 1.170	0.598 / 1.170	ms
Subphase: Concurrent Classes Purge	0.112 / 0.112	0.293 / 0.709	0.293 / 0.709	0.293 / 0.709	ms
Subphase: Concurrent Classes Unlink	0.897 / 0.897	5.852 / 18.071	5.852 / 18.071	5.852 / 18.071	ms
Subphase: Concurrent Mark	315.060 / 317.380	814.912 / 3055.673	814.912 / 3055.673	814.912 / 3055.673	ms
Subphase: Concurrent Mark Idle	1.075 / 1.161	1.958 / 21.000	1.958 / 21.000	1.958 / 21.000	ms
Subphase: Concurrent Mark Try Flush	0.137 / 0.449	6.628 / 51.534	6.628 / 51.534	6.628 / 51.534	ms
...					



=== Garbage Collection Statistics =====

	Last 10s	Last 10m	Last 10h	Total	
	Avg / Max	Avg / Max	Avg / Max	Avg / Max	
Collector: Garbage Collection Cycle	393.373 / 393.373	2457.026 / 4420.944	2457.026 / 4420.944	2457.026 / 4420.944	ms
Contention: Mark Segment Reset Contention	4 / 47	1 / 165	1 / 165	1 / 165	ops/s
Contention: Mark SeqNum Reset Contention	0 / 0	0 / 1	0 / 1	0 / 1	ops/s
Contention: Relocation Contention	0 / 6	13 / 1609	13 / 1609	13 / 1609	ops/s
Critical: Allocation Stall	0.000 / 0.000	0.000 / 0.000	0.000 / 0.000	0.000 / 0.000	ms
Critical: Allocation Stall	0 / 0	0 / 0	0 / 0	0 / 0	ops/s
Critical: GC Locker Stall	0.000 / 0.000	0.000 / 0.000	0.000 / 0.000	0.000 / 0.000	ms
Critical: GC Locker Stall	0 / 0	0 / 0	0 / 0	0 / 0	ops/s

	Last 10s	Last 10m	Last 10h	Total	
	Avg / Max	Avg / Max	Avg / Max	Avg / Max	
Pause Mark End	0.063 / 0.063	0.411 / 1.142	0.411 / 1.142	0.411 / 1.142	ms
Pause Mark Start	0.250 / 0.250	0.407 / 0.706	0.407 / 0.706	0.407 / 0.706	ms
Pause Relocate Start	0.513 / 0.513	0.598 / 1.170	0.598 / 1.170	0.598 / 1.170	ms

Memory: Page Cache Hit LS	2 / 7	127 / 2928	127 / 2928	127 / 2928	ops/s
Memory: Page Cache Miss	0 / 0	0 / 0	0 / 0	0 / 0	ops/s
Memory: Uncommit	0 / 0	0 / 0	0 / 0	0 / 0	MB/s
Memory: Undo Object Allocation Failed	0 / 0	0 / 53	0 / 53	0 / 53	ops/s
Memory: Undo Object Allocation Succeeded	0 / 6	12 / 1609	12 / 1609	12 / 1609	ops/s
Memory: Undo Page Allocation	0 / 0	6 / 463	6 / 463	6 / 463	ops/s
Phase: Concurrent Mark	317.841 / 317.841	1564.321 / 3056.192	1564.321 / 3056.192	1564.321 / 3056.192	ms
Phase: Concurrent Mark Continue	0.000 / 0.000	8.173 / 8.173	8.173 / 8.173	8.173 / 8.173	ms
Phase: Concurrent Process Non-Strong References	1.533 / 1.533	16.783 / 57.370	16.783 / 57.370	1783 / 57.370	ms
Phase: Concurrent Relocate	49.909 / 49.909	748.800 / 1299.535	748.800 / 1299.535	748.800 / 1299.535	ms
Phase: Concurrent Reset Relocation Set	5.138 / 5.138	4.566 / 12.218	4.566 / 12.218	4.566 / 12.218	ms
Phase: Concurrent Select Relocation Set	16.672 / 16.672	111.356 / 259.445	111.356 / 259.445	111.356 / 259.445	ms
Phase: Pause Mark End	0.063 / 0.063	0.411 / 1.142	0.411 / 1.142	0.411 / 1.142	ms
Phase: Pause Mark Start	0.250 / 0.250	0.407 / 0.706	0.407 / 0.706	0.407 / 0.706	ms
Phase: Pause Relocate Start	0.513 / 0.513	0.598 / 1.170	0.598 / 1.170	0.598 / 1.170	ms
Subphase: Concurrent Classes Purge	0.112 / 0.112	0.293 / 0.709	0.293 / 0.709	0.293 / 0.709	ms
Subphase: Concurrent Classes Unlink	0.897 / 0.897	5.852 / 18.071	5.852 / 18.071	5.852 / 18.071	ms
Subphase: Concurrent Mark	315.060 / 317.380	814.912 / 3055.673	814.912 / 3055.673	814.912 / 3055.673	ms
Subphase: Concurrent Mark Idle	1.075 / 1.161	1.958 / 21.000	1.958 / 21.000	1.958 / 21.000	ms
Subphase: Concurrent Mark Try Flush	0.137 / 0.449	6.628 / 51.534	6.628 / 51.534	6.628 / 51.534	ms

...



=== Garbage Collection Statistics =====					
	Last 10s	Last 10m	Last 10h	Total	
	Avg / Max	Avg / Max	Avg / Max	Avg / Max	
Collector: Garbage Collection Cycle	393.373 / 393.373	2457.026 / 4420.944	2457.026 / 4420.944	2457.026 / 4420.944	ms
Contention: Mark Segment Reset Contention	4 / 47	1 / 165	1 / 165	1 / 165	ops/s
Contention: Mark SeqNum Reset Contention	0 / 0	0 / 1	0 / 1	0 / 1	ops/s
Contention: Relocation Contention	0 / 6	13 / 1609	13 / 1609	13 / 1609	ops/s
Critical: Allocation Stall	0.000 / 0.000	0.000 / 0.000	0.000 / 0.000	0.000 / 0.000	ms
Critical: Allocation Stall	0 / 0	0 / 0	0 / 0	0 / 0	ops/s
Critical: GC Locker Stall	0.000 / 0.000	0.000 / 0.000	0.000 / 0.000	0.000 / 0.000	ms
Critical: GC Locker Stall	0 / 0	0 / 0	0 / 0	0 / 0	ops/s
Memory: Allocation Rate	862 / 2892	2229 / 13172	2229 / 13172	2229 / 13172	MB/s
Memory: Heap Used After Mark	6038 / 6038	55657 / 107472	55657 / 107472	55657 / 107472	MB
Memory: Heap Used After Relocation	1496 / 1496	14059 / 31084	14059 / 31084	14059 / 31084	MB
Memory: Heap Used Before Mark	6038 / 6038	47444 / 94452	47444 / 94452	47444 / 94452	MB
Memory: Heap Used Before Relocation	3452 / 3452	47152 / 94430	47152 / 94430	47152 / 94430	MB
Memory: Out Of Memory	0 / 0	0 / 0	0 / 0	0 / 0	ops/s
Memory: Page Cache Flush	0 / 0	0 / 0	0 / 0	0 / 0	MB/s
Memory: Page Cache Hit L1	350 / 1183	765 / 3549	765 / 3549	765 / 3549	ops/s
Memory: Page Cache Hit L2	0 / 0	4 / 234	4 / 234	4 / 234	ops/s
Memory: Page Cache Hit L3	2 / 7	127 / 2928	127 / 2928	127 / 2928	ops/s
Memory: Page Cache Miss	0 / 0	0 / 0	0 / 0	0 / 0	ops/s
Memory: Uncommit	0 / 0	0 / 0	0 / 0	0 / 0	MB/s
Memory: Undo Object Allocation Failed	0 / 0	0 / 53	0 / 53	0 / 53	ops/s
Memory: Undo Object Allocation Succeeded	0 / 6	12 / 1609	12 / 1609	12 / 1609	ops/s
Memory: Undo Page Allocation	0 / 0	6 / 463	6 / 463	6 / 463	ops/s
Phase: Concurrent Mark	317.841 / 317.841	1564.321 / 3056.192	1564.321 / 3056.192	1564.321 / 3056.192	ms
Phase: Concurrent Mark Continue	0.000 / 0.000	8.173 / 8.173	8.173 / 8.173	8.173 / 8.173	ms
Phase: Concurrent Process Non-Strong References	1.533 / 1.533	16.783 / 57.370	16.783 / 57.370	1783 / 57.370	ms
Phase: Concurrent Relocate	49.909 / 49.909	748.800 / 1299.535	748.800 / 1299.535	748.800 / 1299.535	ms
Phase: Concurrent Reset Relocation Set	5.138 / 5.138	4.566 / 12.218	4.566 / 12.218	4.566 / 12.218	ms
Phase: Concurrent Select Relocation Set	16.672 / 16.672	111.356 / 259.445	111.356 / 259.445	111.356 / 259.445	ms
Phase: Pause Mark End	0.063 / 0.063	0.411 / 1.142	0.411 / 1.142	0.411 / 1.142	ms
Phase: Pause Mark Start	0.250 / 0.250	0.407 / 0.706	0.407 / 0.706	0.407 / 0.706	ms
Phase: Pause Relocate Start	0.513 / 0.513	0.598 / 1.170	0.598 / 1.170	0.598 / 1.170	ms
Subphase: Concurrent Classes Purge	0.112 / 0.112	0.293 / 0.709	0.293 / 0.709	0.293 / 0.709	ms
Subphase: Concurrent Classes Unlink	0.897 / 0.897	5.852 / 18.071	5.852 / 18.071	5.852 / 18.071	ms
Subphase: Concurrent Mark	315.060 / 317.380	814.912 / 3055.673	814.912 / 3055.673	814.912 / 3055.673	ms
Subphase: Concurrent Mark Idle	1.075 / 1.161	1.958 / 21.000	1.958 / 21.000	1.958 / 21.000	ms
Subphase: Concurrent Mark Try Flush	0.137 / 0.449	6.628 / 51.534	6.628 / 51.534	6.628 / 51.534	ms
...					





# Roadmap



Sub-millisecond max pause times

Generational ZGC

# Sub-millisecond max pause times

# Initial Goals

Max GC pause time

**10<sub>ms</sub>**

Multi-terabyte heaps

**TB**

Max throughput reduction

**15%**



# Current Goals

Max GC pause time

~~10<sub>ms</sub>~~  
1<sub>ms</sub>

Multi-terabyte heaps

TB

Max throughput reduction

15%

# What's Concurrent?

\*) Old Gen Only  
 \*\*) Not yet

	Serial	Parallel	G1	CMS	ZGC
Marking	-	-	✓ *	✓ *	✓
Compaction	-	-	-	-	✓
Reference Processing	-	-	-	-	✓
Relocation Set Selection	-	-	-	-	✓
StringTable Cleaning	-	-	-	-	✓
JNI WeakRef Cleaning	-	-	-	-	✓
JNI GlobalRefs Scanning	-	-	-	-	✓
Class Unloading	-	-	-	-	✓
Thread Stack Scanning	-	-	-	-	- **

\*) Old Gen Only

# What's Concurrent?

	Serial	Parallel	G1	CMS	ZGC
Marking	-	-	✓ *	✓ *	✓
Compaction	-	-	-	-	✓
Reference Processing	-	-	-	-	✓
Relocation Set Selection	-	-	-	-	✓
StringTable Cleaning	-	-	-	-	✓
JNI WeakRef Cleaning	-	-	-	-	✓
JNI GlobalRefs Scanning	-	-	-	-	✓
Class Unloading	-	-	-	-	✓
Thread Stack Scanning	-	-	-	-	✓



ZGC pause times do increase  
with the root-set size

(Number of Java Threads)



do not  
ZGC pause times ~~do~~ increase  
with the root-set size

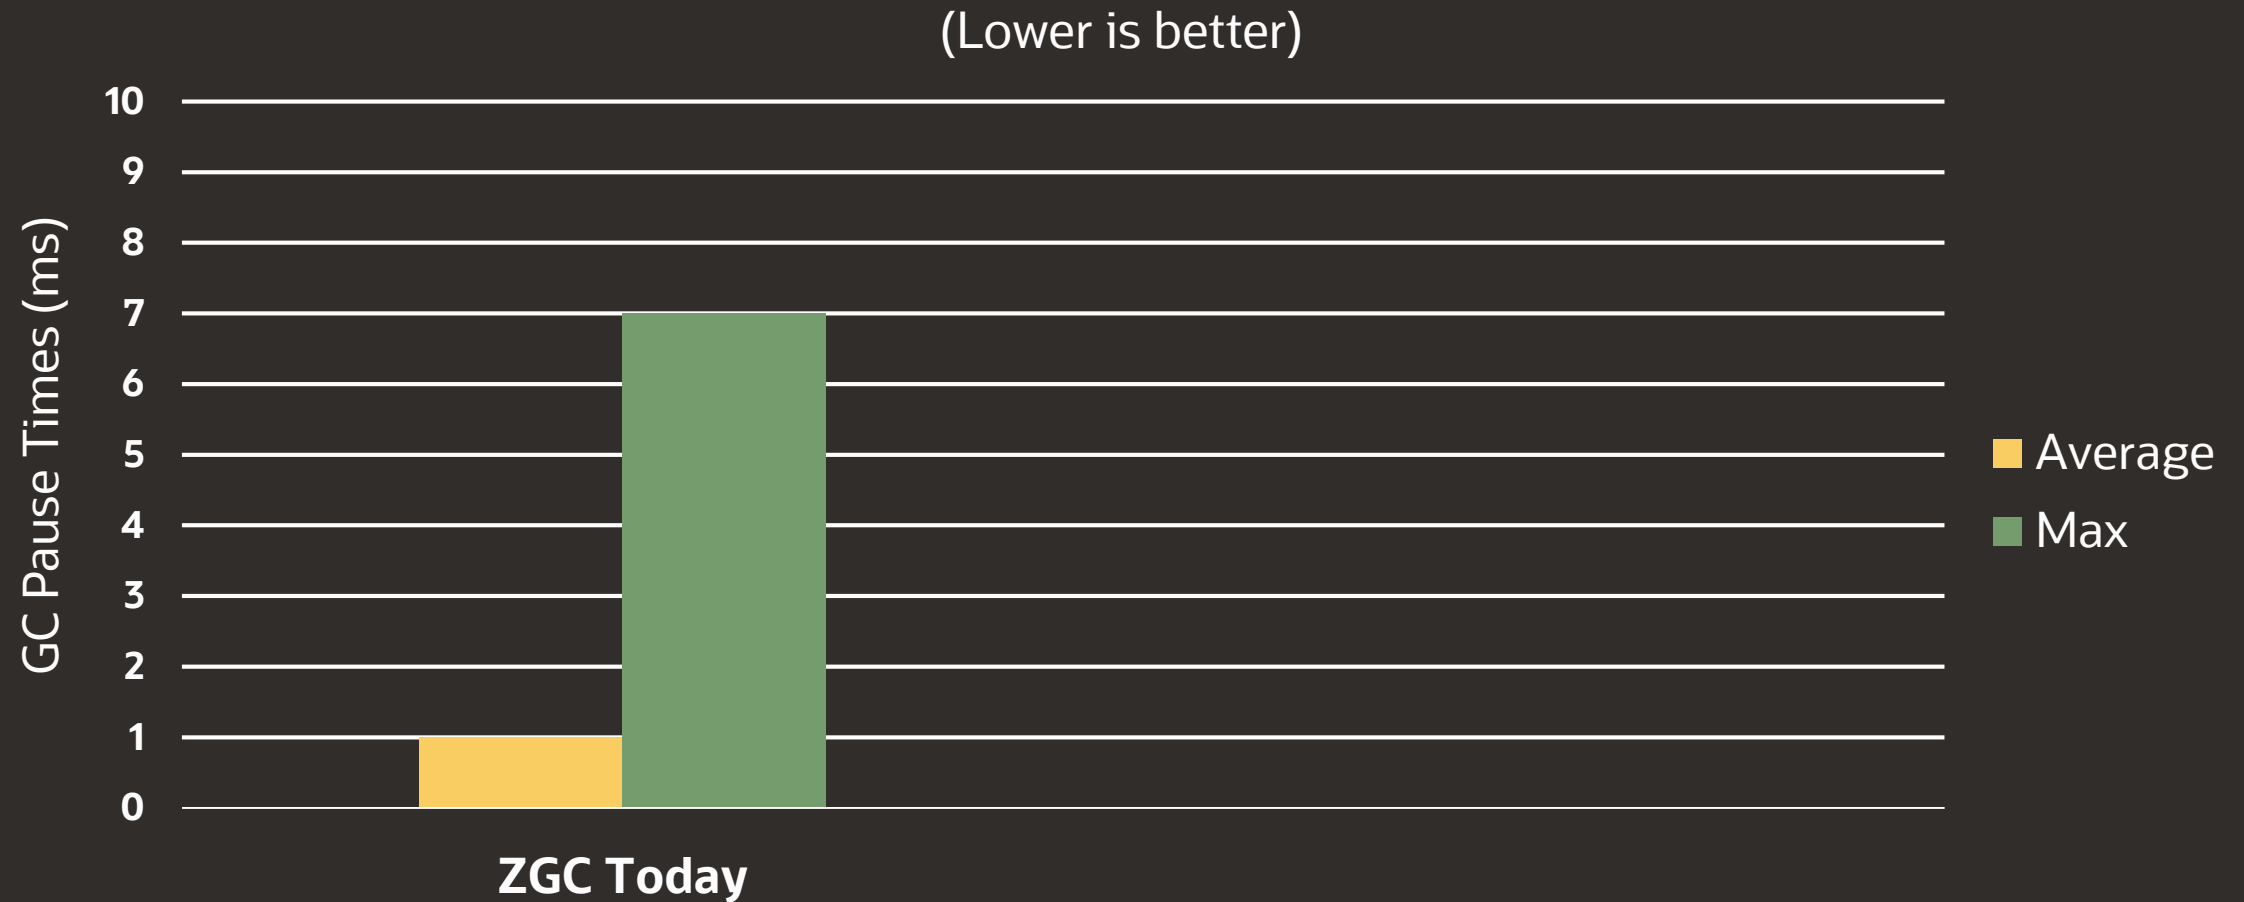
(Number of Java Threads)



# JEP 376: ZGC: Concurrent Thread-Stack Processing

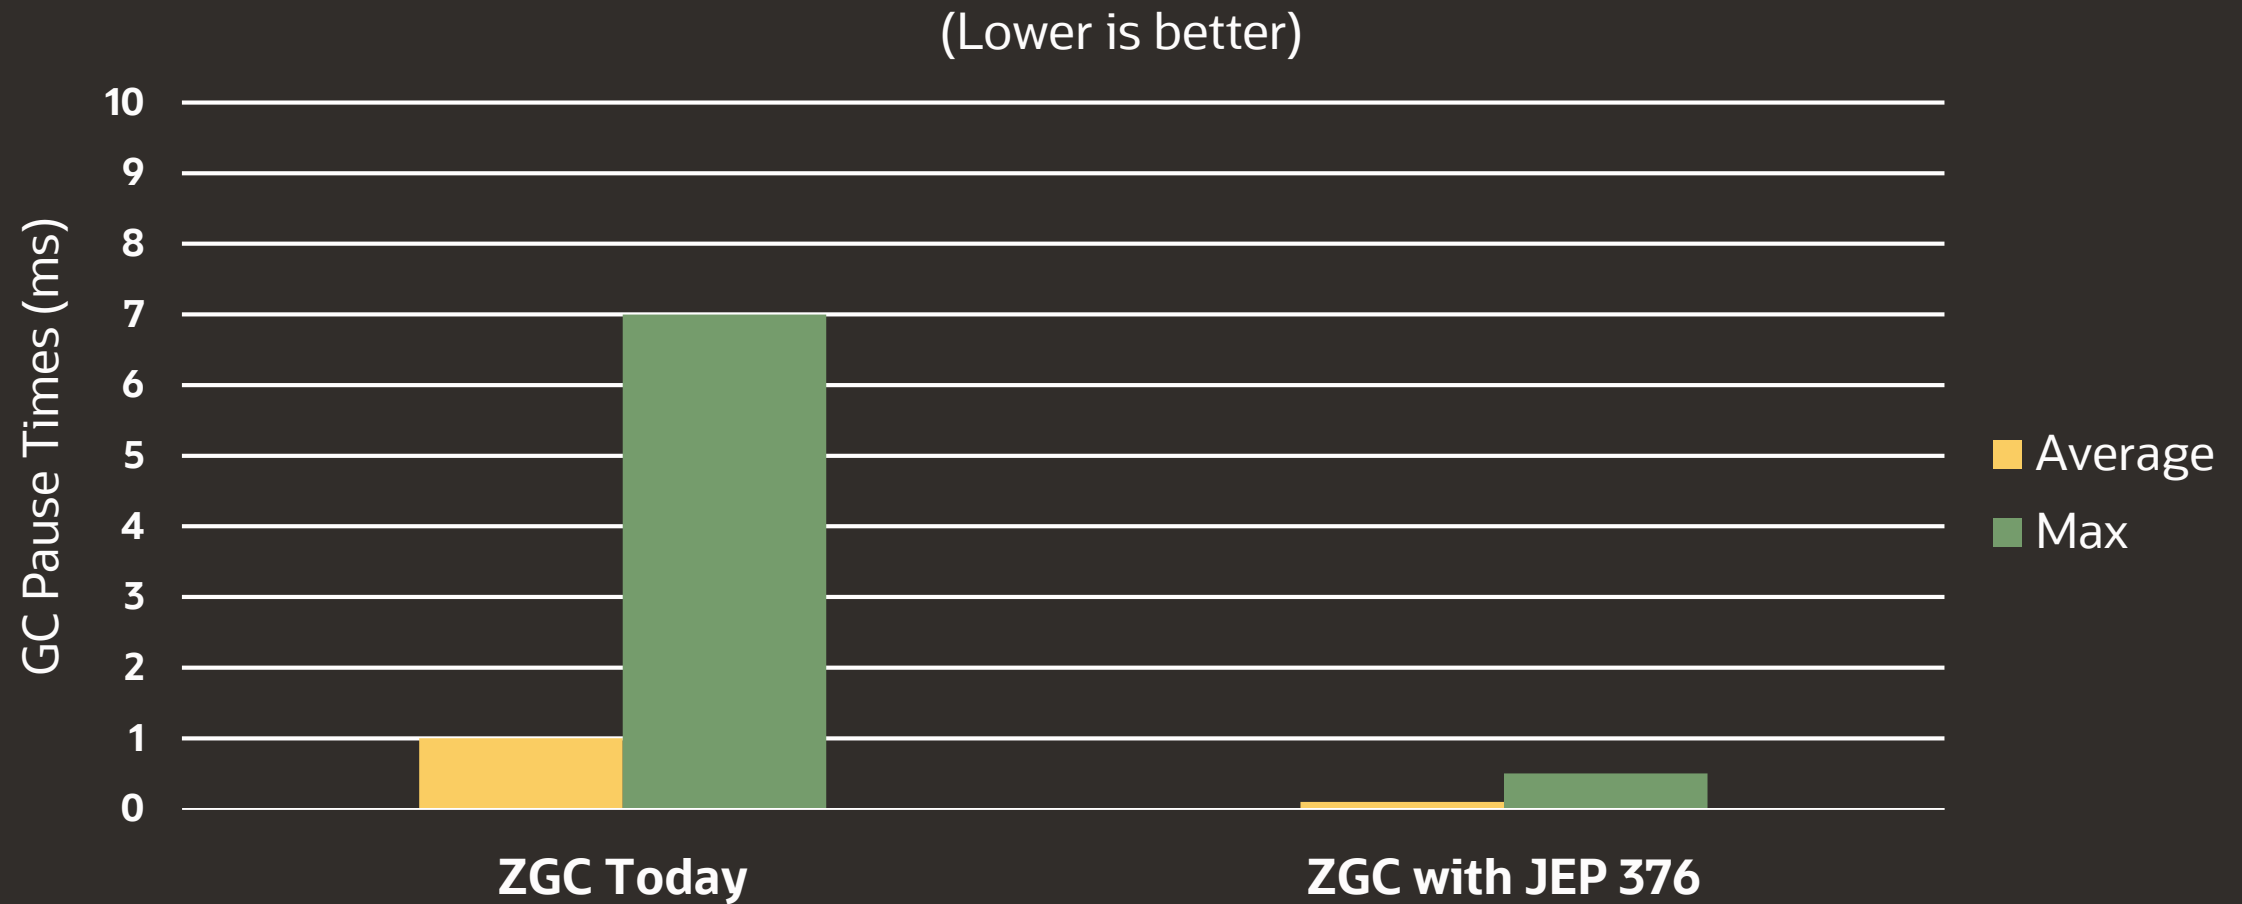
# SPECjbb2015 – Preliminary Results

**3T** Heap  
**224** Hyper-threads (Intel)  
**~2100** Java threads



# SPECjbb2015 – Preliminary Results

3T Heap  
224 Hyper-threads (Intel)  
~2100 Java threads







# Generational ZGC

# Generational ZGC



## Young/Old Generation

- Exploit the fact that most objects are short-lived

## Benefits

- Withstand higher allocation rates
- Lower heap headroom
- Lower CPU usage

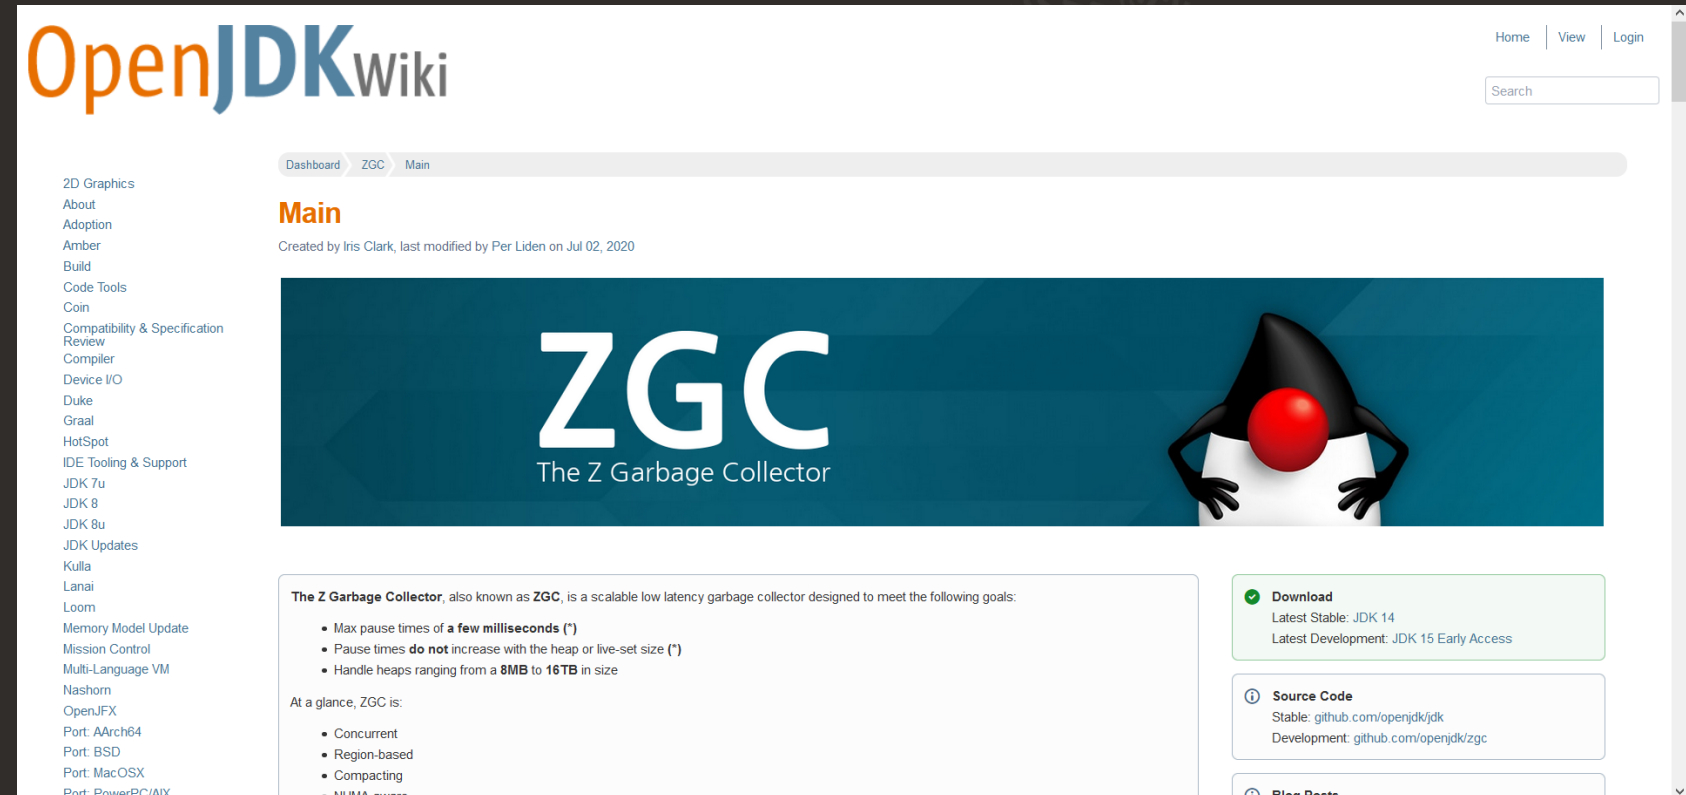
## Work in Progress



# More Information

# More Information

## ZGC Wiki



The screenshot shows the OpenJDK Wiki page for ZGC. The page has a dark blue header with the OpenJDK Wiki logo and navigation links (Home, View, Login). A search bar is in the top right. The main content area has a breadcrumb trail (Dashboard > ZGC > Main) and a title 'Main' with a subtitle 'Created by Iris Clark, last modified by Per Liden on Jul 02, 2020'. A large banner features the text 'ZGC The Z Garbage Collector' and a cartoon penguin character. Below the banner, a box titled 'The Z Garbage Collector' lists its goals: max pause times of a few milliseconds, pause times that do not increase with heap size, and handling heaps from 8MB to 16TB. It also lists features like Concurrent, Region-based, Compacting, and NUMA aware. On the right, there are sections for 'Download' (Latest Stable: JDK 14, Latest Development: JDK 15 Early Access), 'Source Code' (Stable: github.com/openjdk/jdk, Development: github.com/openjdk/zgc), and 'Blog Posts'.

OpenJDK Wiki

Home | View | Login

Search

Dashboard > ZGC > Main

### Main

Created by Iris Clark, last modified by Per Liden on Jul 02, 2020

# ZGC

The Z Garbage Collector

**The Z Garbage Collector**, also known as **ZGC**, is a scalable low latency garbage collector designed to meet the following goals:

- Max pause times of **a few milliseconds** (\*)
- Pause times **do not** increase with the heap or live-set size (\*)
- Handle heaps ranging from a **8MB** to **16TB** in size

At a glance, ZGC is:

- Concurrent
- Region-based
- Compacting
- NUMA aware

**Download**  
Latest Stable: JDK 14  
Latest Development: JDK 15 Early Access

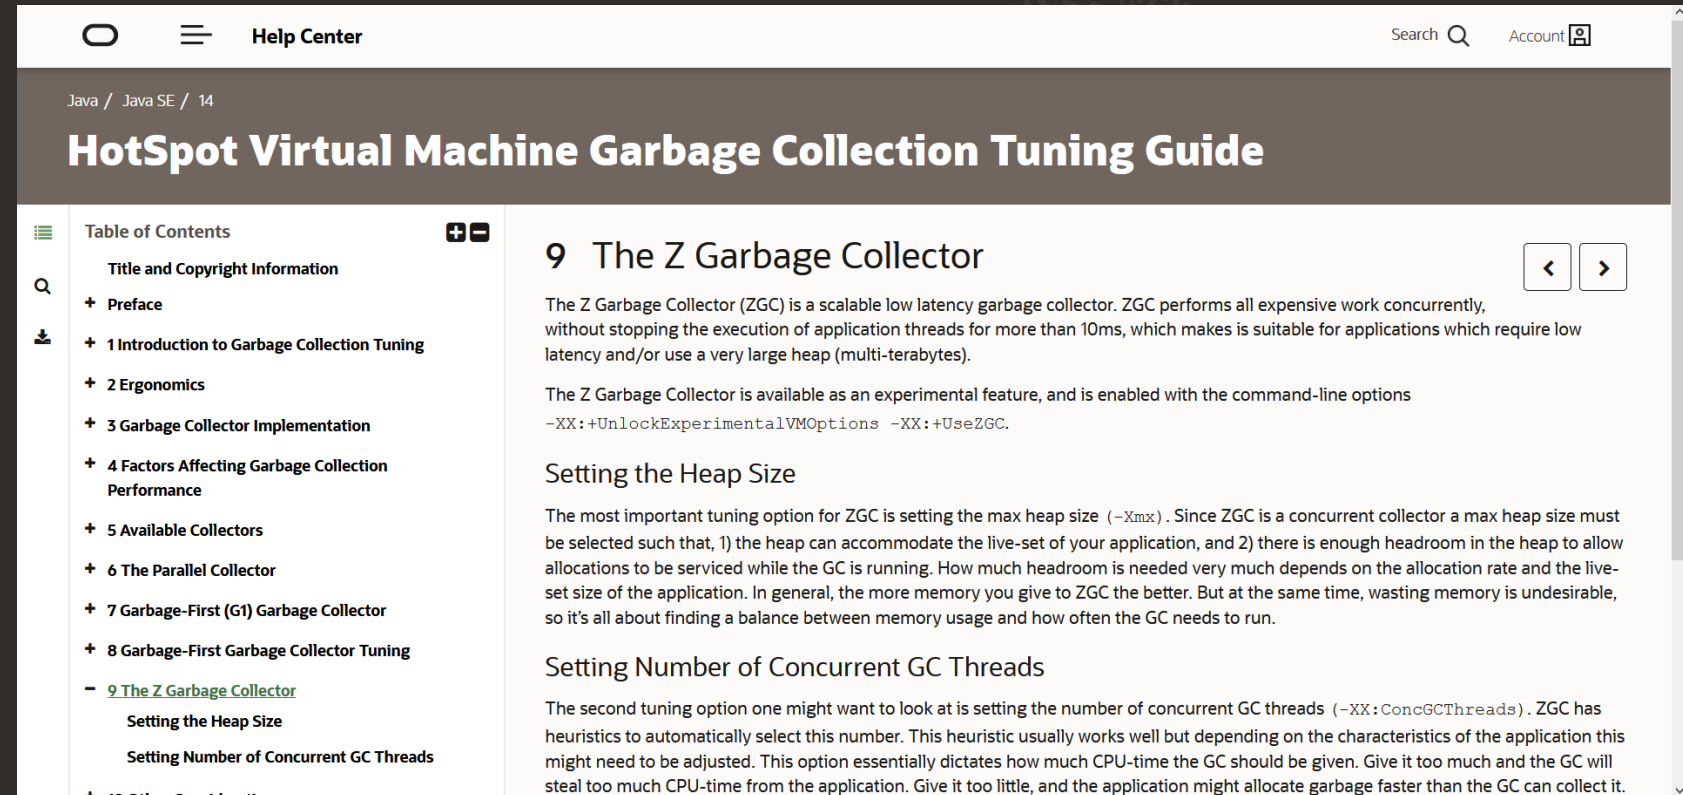
**Source Code**  
Stable: [github.com/openjdk/jdk](https://github.com/openjdk/jdk)  
Development: [github.com/openjdk/zgc](https://github.com/openjdk/zgc)

**Blog Posts**

<https://wiki.openjdk.java.net/display/zgc>

# More Information

## GC Tuning Guide



The screenshot shows the Oracle Help Center interface. At the top, there's a navigation bar with the Oracle logo, a menu icon, and the text "Help Center". On the right, there are search and account icons. Below the navigation bar, the breadcrumb "Java / Java SE / 14" is visible. The main heading is "HotSpot Virtual Machine Garbage Collection Tuning Guide". On the left, a table of contents lists various topics, with "9 The Z Garbage Collector" highlighted. The main content area displays the title "9 The Z Garbage Collector" followed by two paragraphs of text. The first paragraph describes the Z Garbage Collector (ZGC) as a scalable low latency garbage collector. The second paragraph states that ZGC is available as an experimental feature and provides command-line options to enable it. Below the text, there are two sub-sections: "Setting the Heap Size" and "Setting Number of Concurrent GC Threads", each with a paragraph of text explaining the tuning options.

Help Center

Search Account

Java / Java SE / 14

## HotSpot Virtual Machine Garbage Collection Tuning Guide

### 9 The Z Garbage Collector

The Z Garbage Collector (ZGC) is a scalable low latency garbage collector. ZGC performs all expensive work concurrently, without stopping the execution of application threads for more than 10ms, which makes it suitable for applications which require low latency and/or use a very large heap (multi-terabytes).

The Z Garbage Collector is available as an experimental feature, and is enabled with the command-line options `-XX:+UnlockExperimentalVMOptions -XX:+UseZGC`.

#### Setting the Heap Size

The most important tuning option for ZGC is setting the max heap size (`-Xmx`). Since ZGC is a concurrent collector a max heap size must be selected such that, 1) the heap can accommodate the live-set of your application, and 2) there is enough headroom in the heap to allow allocations to be serviced while the GC is running. How much headroom is needed very much depends on the allocation rate and the live-set size of the application. In general, the more memory you give to ZGC the better. But at the same time, wasting memory is undesirable, so it's all about finding a balance between memory usage and how often the GC needs to run.

#### Setting Number of Concurrent GC Threads

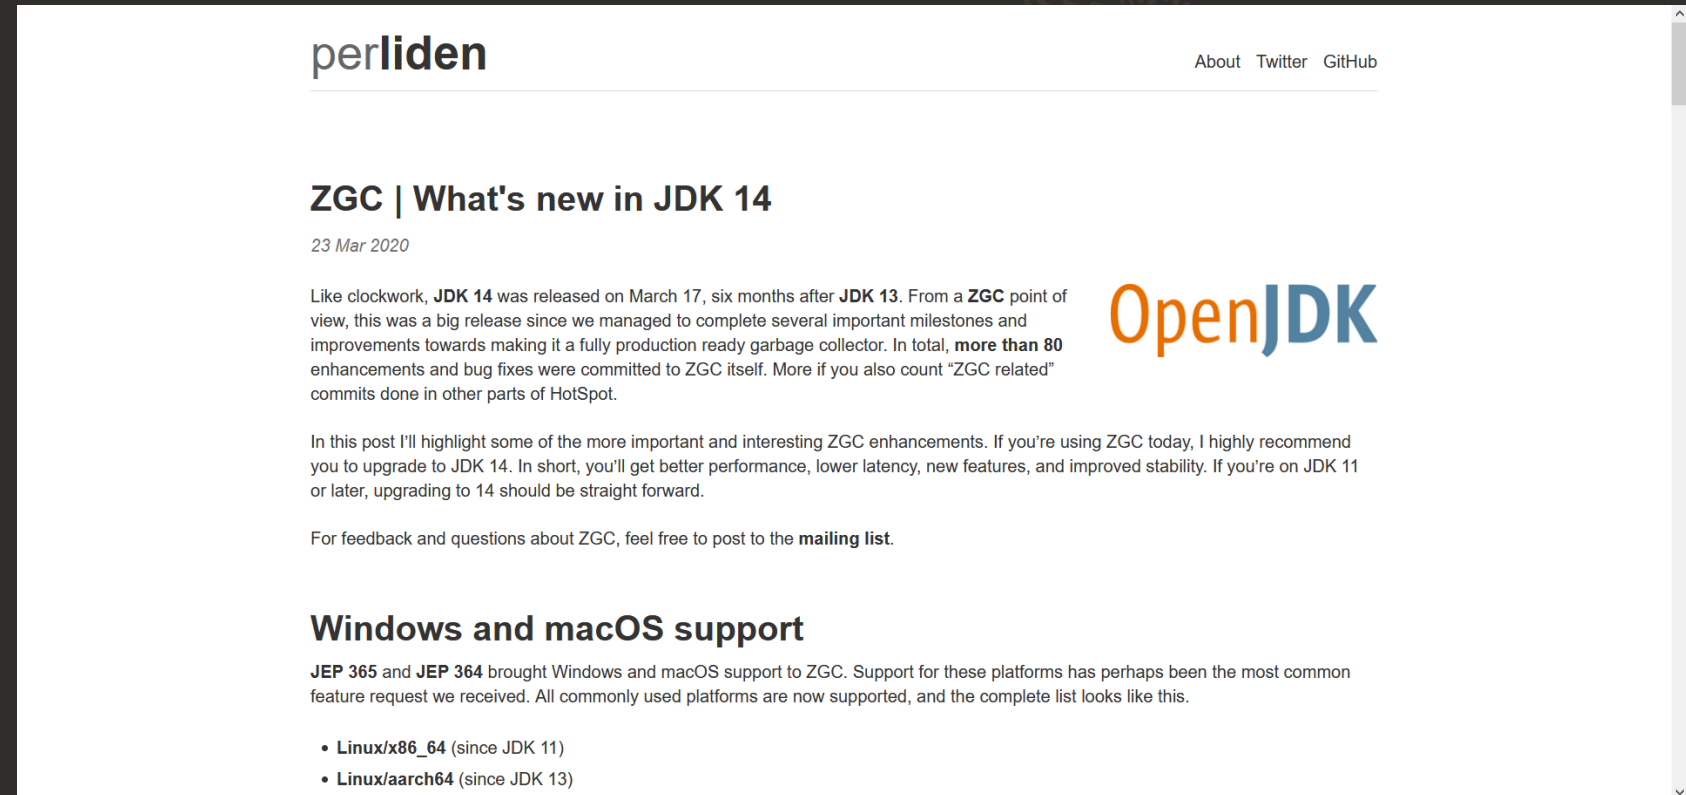
The second tuning option one might want to look at is setting the number of concurrent GC threads (`-XX:ConcGCThreads`). ZGC has heuristics to automatically select this number. This heuristic usually works well but depending on the characteristics of the application this might need to be adjusted. This option essentially dictates how much CPU-time the GC should be given. Give it too much and the GC will steal too much CPU-time from the application. Give it too little, and the application might allocate garbage faster than the GC can collect it.

<https://docs.oracle.com/en/java/javase/15/gctuning/>



# More Information

## My Blog



<https://mallocc.se>



# Download JDK 15!

<https://jdk.java.net/15>

# Q&A



ORACLE

